

Applied Machine Learning - Basic

Prof. Daniele Bonacorsi

Lecture 3

Data Science and Computation PhD + Master in Bioinformatics
University of Bologna

Communications

Miscellanea

I am protecting today's lecture and Friday's lecture

- Be ready for possible announcements, though → details at the lecture

Recap from **Lecture2**

Model representation

(...)

Univariate linear regression - Gradient Descent

Multivariate linear regression

Gradient Descent

Feature scaling and standardisation/normalisation

Polynomial regression

Analytical solution: Normal equation

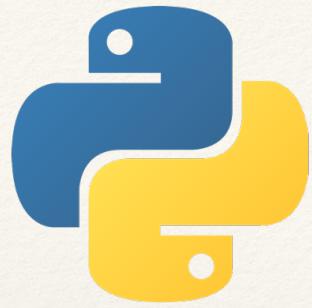
NEXT: material on Py/Jupyter/Colab + Logistic Regression

Applied Machine Learning - Basic

Prof. Daniele Bonacorsi

Hands-on:
Python

Data Science and Computation PhD + Master in Bioinformatics
University of Bologna



Python introduction (part 0)

DISCLAIMER: "This is not a programming course", remember?

Programming languages: the landscape

Let's see some examples

- **high-level** languages
 - ❖ e.g. C, C++, Fortran, Java, ..

Programming languages: the landscape

Let's see some examples

- **high-level** languages
 - ❖ e.g. C, C++, Fortran, Java, ..
- **very high-level interactive** computational environments (general-purpose)
 - ❖ e.g. Matlab/Octave, Mathematica, Python, ..

Programming languages: the landscape

Let's see some examples

- **high-level** languages
 - ❖ e.g. C, C++, Fortran, Java, ..
- **very high-level interactive** computational environments (general-purpose)
 - ❖ e.g. Matlab/Octave, Mathematica, Python, ..
- **text processing** utilities and **scripting** languages
 - ❖ e.g. Unix shell scripting, sed, awk, Perl, ..

Programming languages: the landscape

Let's see some examples

- **high-level** languages
 - ❖ e.g. C, C++, Fortran, Java, ..
- **very high-level interactive** computational environments (general-purpose)
 - ❖ e.g. Matlab/Octave, Mathematica, Python, ..
- **text processing** utilities and **scripting** languages
 - ❖ e.g. Unix shell scripting, sed, awk, Perl, ..
- **special-purpose** languages
 - ❖ e.g. R for stats, LabVIEW for lab activities, SQL for databases, ..

Very-high level languages: Python



General purpose, interpreted programming language

Compact and readable - like pseudo-code

- looks like an English language description of what needs to be done

Widely used for scientific programming (→ also data science)

Versatile to use, and very forgiving

- thus enabling quick coding times

Recognisable (e.g. indentation)

- e.g. Python uses whitespace indentation to delimit code blocks (instead of curly braces, as in C, or keywords as 'end' in Fortran)

Chinese Python

[周蟒 → “Zhōu mǎng”]

<http://zh.wikipedia.org/wiki/ZhPy> →

zh.wikipedia.org/wiki/周蟒

维基百科，自由的百科全书

条目 讨论 不转换 汉 漢 阅读 编辑 查看历史 搜索维基百科 [关闭]

周蟒 [编辑]

维基百科，自由的百科全书
(重定向自ZhPy)

本條目需要补充更多來源。 (2016年11月30日)
请协助添加多方面可靠来源以改善这篇条目，无法查证的内容可能會因為异议提出而移除。

周蟒，又名zphy，是一個輕量的，與Python語言互相兼容的中文Python語言。讓使用者可以使用純中文語句（繁體或簡體）來編寫程式。目前主要適用於教學上。

周蟒中文程式語言目前已不再更新。提供電子書、API、完整測試用例的開放原始碼中文程式語言。

周蟒中文程式語言的目標是協助使用者透過中文程式語言學習程式語言，進而接觸世界上大部份的程式語言，而不是脫離現實世界。周蟒作者gasolin提出了周蟒編程風格^[1]與周蟒中文程式語言的四不一沒有，透過程式語言的約定，使用周蟒寫出的中文程式保有了易於閱讀的特性，並可完全轉換成英文Python語言程式。

周蟒語言擁有多項Python語言的所有特性，如高效率的高階資料結構、簡單而有效的物件導向程式設計方式等等。由於周蟒語言完全相容Python程式語言，所以可以取用所有Python程式語言資源。

周蟒語言的長處是在於發揮“完全相容Python程式語言”的中文程式語言的優點，所有語法，關鍵詞都依照Python語言的風格。學習周蟒語言後要橋接到Python語言相當容易。

周蟒也同時提供咬一口周蟒中文程式語言^[2]電子書，所有範例都一併提供與python語言程式碼的對照。

周蟒是開放原始碼的，可自由下載使用。

周蟒提供編譯器與互動式直譯器，也支援中文腳本執行。

周蟒的Hello World程序 [编辑]

下面是一個在標準輸出設備上輸出Hello World的簡單程式，這種程式通常作為開始學習程式語言時的第一個程式：

```
#!/usr/bin/env zphy
印出 "哈囉，世界"
```

周蟒

编程范型	multi-paradigm
实作者	Fred Lin (gasolin)
发行时间	2007年
预定版本	3.0.0a1 (2011年6月25日)
型態系統	Strong, dynamic ("duck")
作業系統	跨平台
許可證	MIT許可證
網站	GitHub
主要實作產品	zphy (即周蟒)
啟發語言	Python, 中蟒

Chinese Python

範例 [编辑]

以下是程式透過編譯器執行的範例：

```
#!/usr/bin/env zhpy
# 檔名: while.py
數字 = 23
運行 = 真
當 運行:
    猜測 = 整數(輸入('輸入一個數字: '))

    如果 猜測 == 數字:
        印出 '恭喜, 你猜對了.'
        運行 = 假 # 這會讓循環語句結束
    假使 猜測 < 數字:
        印出 '錯了, 數字再大一點.'
    否則:
        印出 '錯了, 數字再小一點.'
    否則:
        印出 '循環語句結束'
    印出 '結束'
```

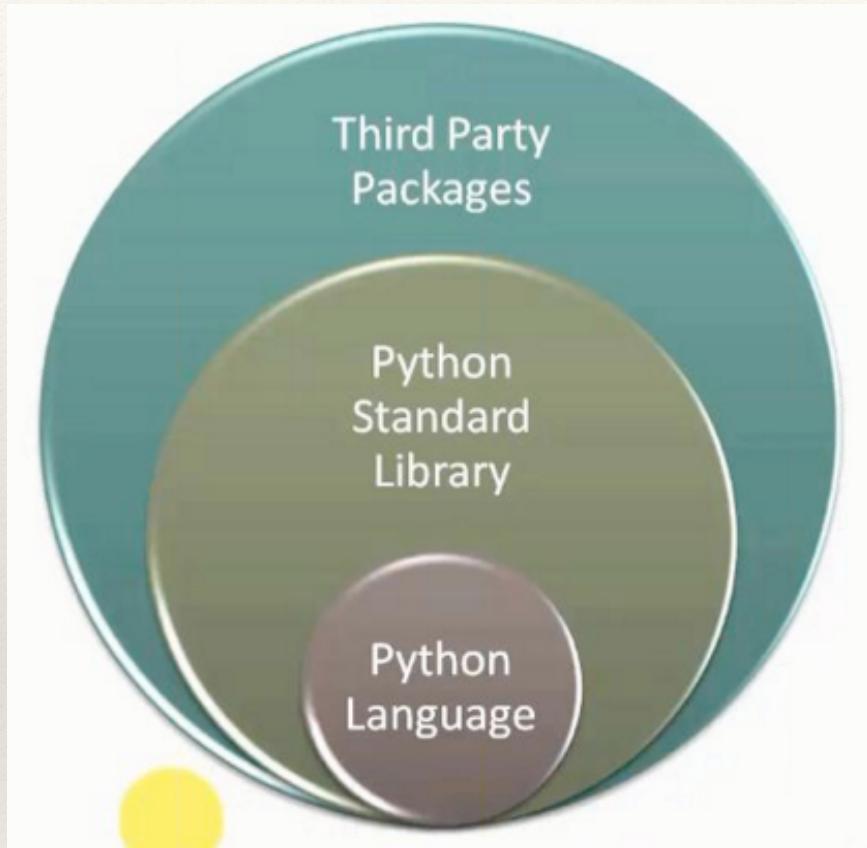
Python 版：

```
#!/usr/bin/env python
# File name: while.twpy
number = 23
running = True
while running:
    guess = int(raw_input('Enter an integer : '))

    if guess == number:
        print 'Congratulations, you guessed it.'
        running = False # this causes the while loop to stop
    elif guess < number:
        print 'No, it is higher than that.'
    else:
        print 'No, it is lower than that.'
else:
    print 'The while loop is over'
print 'Done'
```

e.g. 印出 [Yìn chū] indeed means “print out”

The Python env



a huge collection of
third party packages

very extensive standard library built
around it - every normal capability you
expect from a programming language
can probably be found in this standard
library

relatively small and
compact language:
www.python.org

More technically, on Python

Python supports **dynamic typing**

Object-oriented language from the ground up

- every function and class is actually an object

Includes **functional programming** features

Supports **meta-programming**

- you can use the Python language to extend Python itself

Extensible in C/C++

- Python can be extended by writing new modules in C/C++
- A number of the third party packages that have been contributed to the Python language are actually written in C/C++

Which Python version?

Python 2.x is the past

- end of support: January 2020

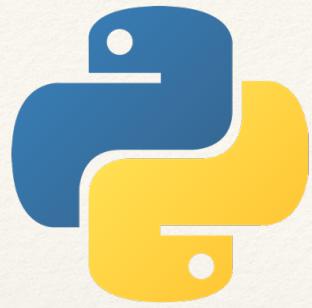
Python 3.x is the present (and the expected future for a while, now)

Unfortunately they are not compatible with each other.

If you start using Python today, you should start with Python 3 **but** you will soon bump into legacy packages written in Python 2

- e.g. machine learning?
- Big Science experiments code base? E.g. many (most) large physics experiments are still at some subversion of Python 2 for stability reasons..

Unfortunately, you might need to be familiar with both 2 and 3, and the differences.. not a big deal, though!



Python light overview (part 1): more on the language environment

DISCLAIMER: "This is not a programming course", remember?

The screenshot shows the Python.org homepage with several UI elements highlighted by red boxes:

- A red box highlights the "PyPI" button in the top navigation bar.
- A red box highlights the "Download" button in the bottom navigation bar.
- A red box highlights the "Docs" button in the bottom navigation bar.

The page content includes:

- The Python logo and the word "python™".
- A search bar with a magnifying glass icon and a "GO" button.
- A menu bar with links: About, Downloads, Documentation, Community, Success Stories, News, Events.
- A code snippet demonstrating Python list comprehensions and the enumerate function.
- A section titled "Compound Data Types" with a brief description and a link to "More about lists in Python 3".
- Pagination numbers 1, 2, 3, 4, 5.
- A central message: "Python is a programming language that lets you work quickly and integrate systems more effectively. [»» Learn More](#)".

Get Started

Whether you're new to programming or an experienced developer, it's easy to learn and use Python.

[Start with our Beginner's Guide](#)

Download

Python source code and installers are available for download for all versions!

Latest: Python 3.10.2

Docs

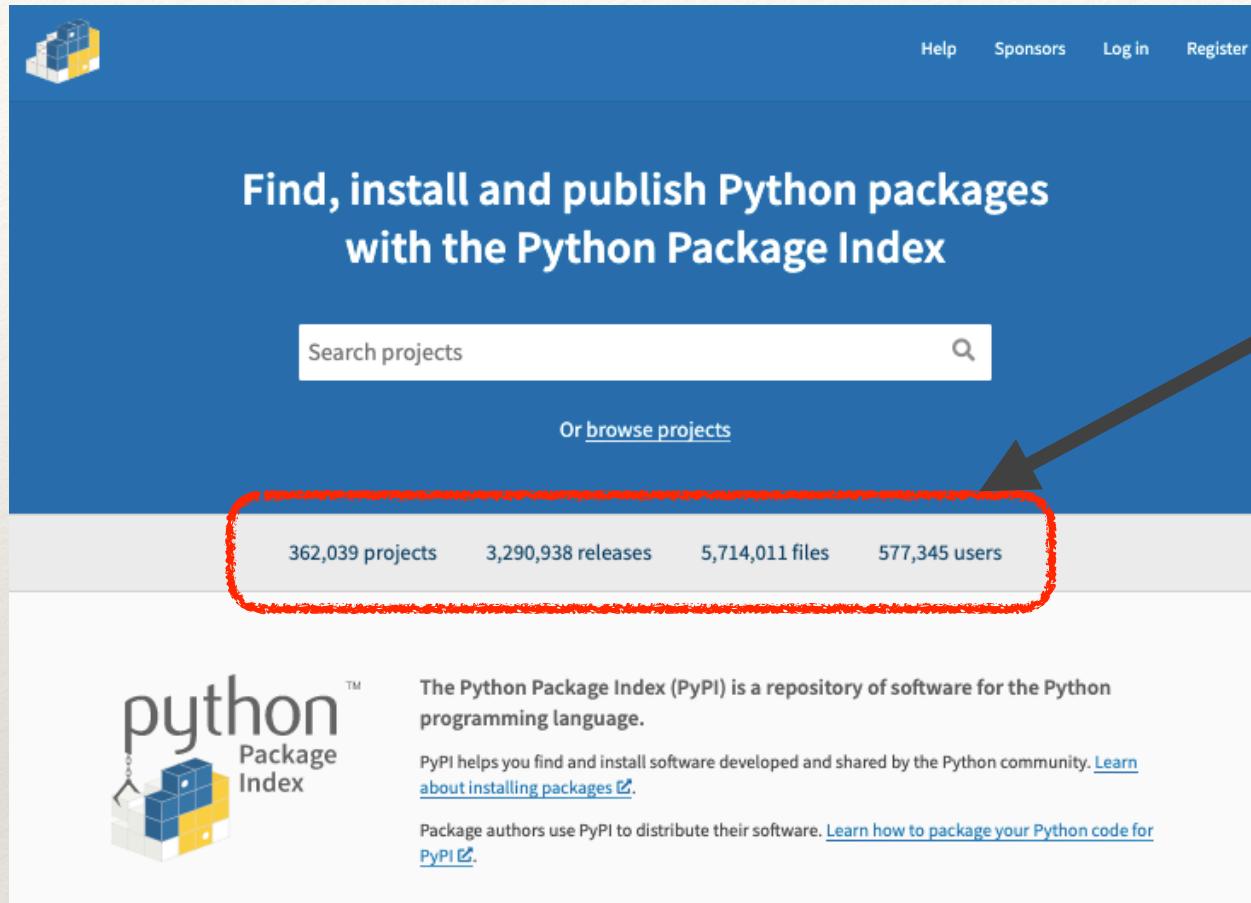
Documentation for Python's standard library, along with tutorials and guides, are available online.

docs.python.org

Jobs

Looking for work or have a Python related position that you're trying to hire for? Our [relaunched community-run job board](#) is the place to go.

Third party packages - pypi.python.org



[snapshot as of March 2022 - numbers may vary]

[last yr, same period of now, it was 3'995k releases...]

You can find them all in the **Python Package Index (PyPI)**, a repository of software for the Python programming language.

- any feature you think you might need in Python, it is probably in the **Standard Library**
- any **application or general purpose utility** you think you might need, it is probably in **PyPI**

pip as a package manager

You will most probably use **pip** to install packages

- google it..
- .. and get familiar with hands-on installations..

Python implementations

C_hython

- the Python reference implementation (from python.org)

Many alternatives... (listed also at python.org itself)

Also a large number of integrated development environment for Python available

- both free and paid for
- same for Python editors

Very popular implementations and/or distributions:

- IPython -> **Jupyter Notebook** (also supports other languages)
- the **Anaconda** distribution: a single distribution with Python, Jupyter and also **Numpy**, **Scipy**, and [many more packages](#) popular for data science applications

Jupyter Notebook

<http://jupyter.org/>



A screenshot of the Jupyter Notebook interface. On the left, there's a sidebar with a "Welcome to the Jupyter Notebook" message and some configuration options. The main area shows a notebook cell titled "Exploring the Lorenz System". The cell contains text about the Lorenz system, three equations, and a description of its chaotic behavior. Below the text is a code cell with Python code for generating a plot. A slider interface allows users to adjust parameters like angle, max_time, sigma, beta, and rho. To the right of the slider is a 3D plot of the Lorenz attractor, showing its characteristic butterfly shape. The plot is composed of many colored lines forming the complex, self-repeating pattern of the system's trajectories.

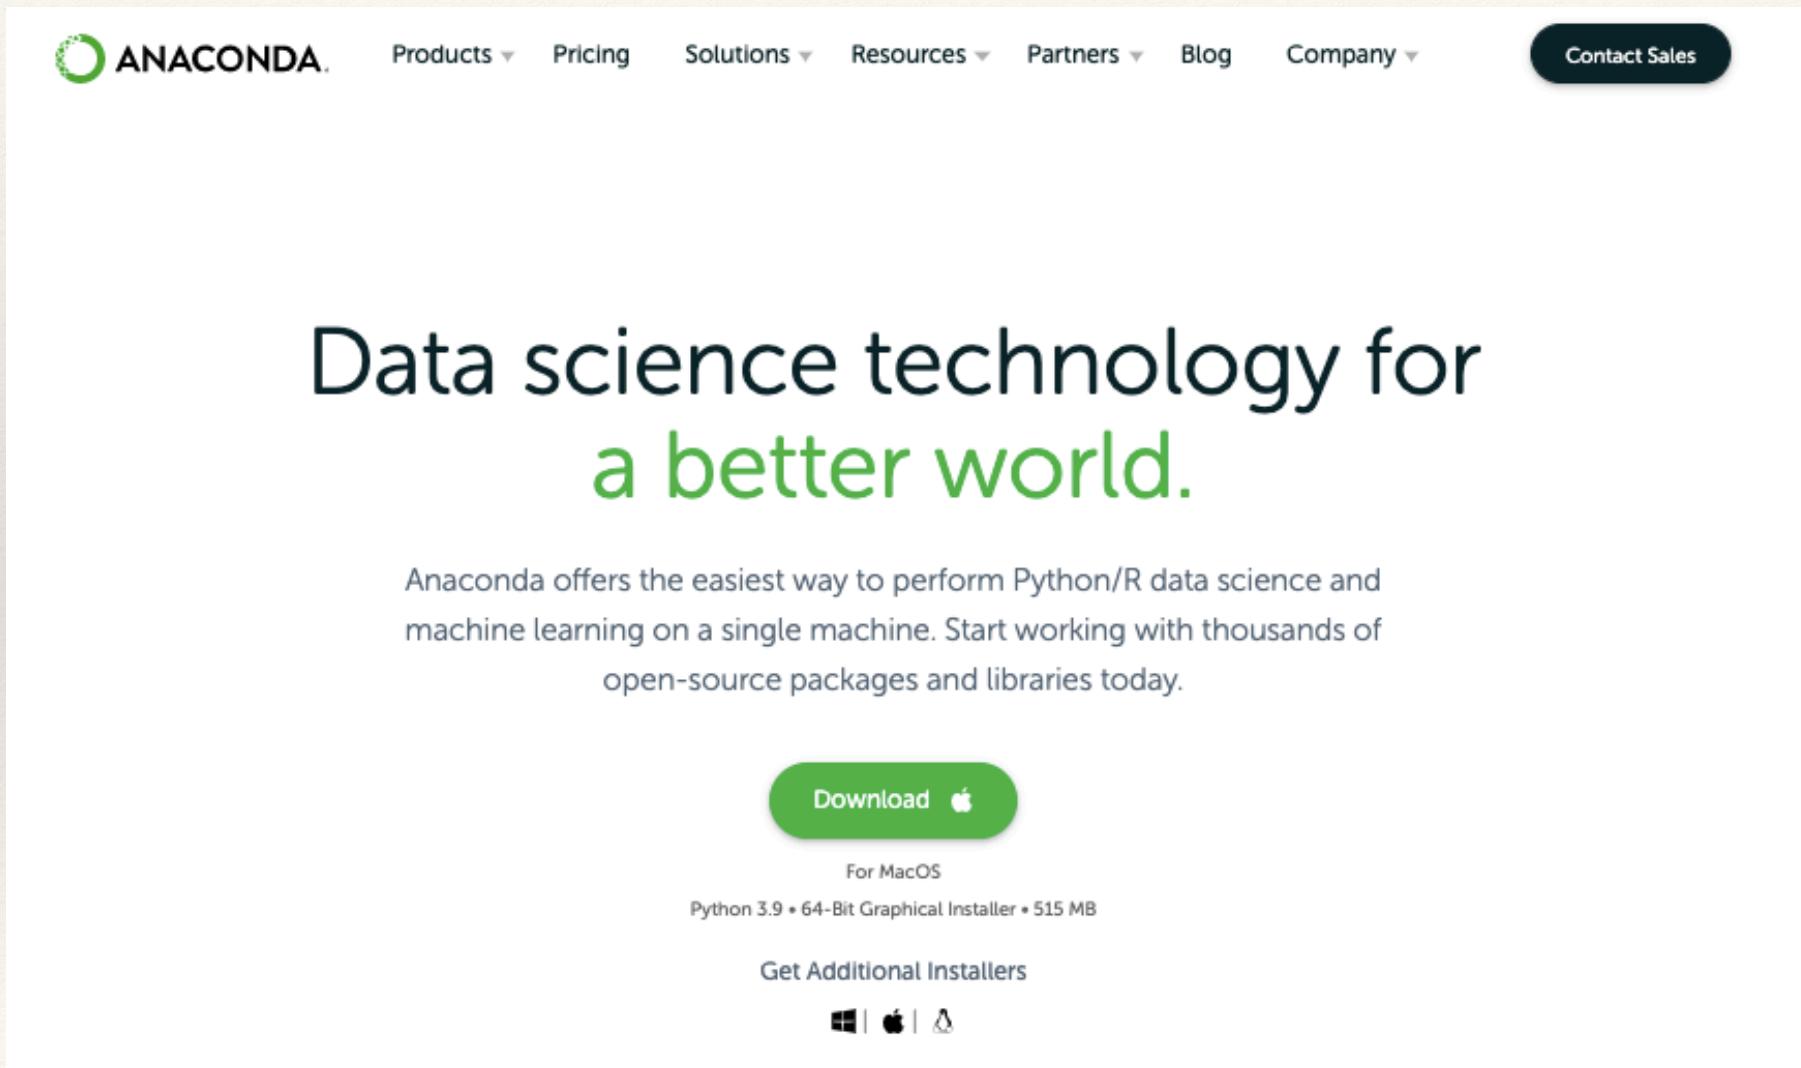
The Jupyter Notebook

The Jupyter Notebook is an open-source web application that allows you to create and share documents that contain live code, equations, visualizations and narrative text. Uses include: data cleaning and transformation, numerical simulation, statistical modeling, data visualization, machine learning, and much more.

More later..

Anaconda

<https://www.anaconda.com/>



The image shows a screenshot of the Anaconda website. At the top, there is a navigation bar with the Anaconda logo, followed by links for Products, Pricing, Solutions, Resources, Partners, Blog, Company, and a Contact Sales button. The main headline reads "Data science technology for a better world." Below the headline, a paragraph describes Anaconda's offerings: "Anaconda offers the easiest way to perform Python/R data science and machine learning on a single machine. Start working with thousands of open-source packages and libraries today." A large green "Download" button with an Apple icon is centered below this text. Below the button, it says "For Mac OS" and provides a download link: "Python 3.9 • 64-Bit Graphical Installer • 515 MB". There is also a "Get Additional Installers" link and icons for Windows, Mac, and Linux.

ANACONDA

Products ▾ Pricing Solutions ▾ Resources ▾ Partners ▾ Blog Company ▾ Contact Sales

Data science technology for a better world.

Anaconda offers the easiest way to perform Python/R data science and machine learning on a single machine. Start working with thousands of open-source packages and libraries today.

Download 

For Mac OS

Python 3.9 • 64-Bit Graphical Installer • 515 MB

Get Additional Installers

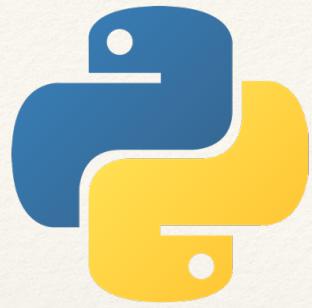
 |  | 

More on Python implementations

<https://www.python.org/download/alternatives/>

<https://wiki.python.org/moin/IntegratedDevelopmentEnvironments>

<https://wiki.python.org/moin/PythonEditors>



Python light overview (part 2): language syntax

DISCLAIMER: "This is not a programming course", remember?

Variable types

1 22 333

Integer

3.14159 1e-6

Floating point

False **True**

Boolean

'Hello' "Let's go!"

String

Comments

```
# This is a regular Python comment  
a = 1    # End-of-line comment  
  
'''A comment can also take the form of a string'''  
  
"""  
A multi-line comment  
"""  
  
...  
Another  
multi-line  
comment  
'''
```

Doc strings

There is a number of third-party packages for interpreting these multiline comments as documentation!

Control statements

As any other language

```
if a == 0:  
    b = 1  
    if c < 0 and d > 0:  
        b = 2  
elif a > 0:  
    b = 3  
else:  
    for i in range(4):  
        b = 4  
        while b > 0:  
            print(b)  
            b = b - 1  
        if i == j: break
```

Note:
INDENTATION is a vital part of the Python syntax as control structure. No “end” keyword. It looks quirky at the beginning, but you will get used to (and will love) it as a neat choice.

Operators

< > == >= <= !=

Comparison

and or not

Boolean

+ - * / // % **

Arithmetic

~ & | >> <<

Bitwise

= += -= *= /= //=% **=

Assignment

in not in

Membership

is is not

Identity

Familiar, if you
know e.g. C..

Range

```
for i in range(4):  
    print(i, end=' ')
```

0 1 2 3

Keyword argument

```
for i in range(2, 6):  
    print(i, end=' ')
```

2 3 4 5

range(Start , Stop , Step)

Functions

Just 2 keywords: **def** and **return**

```
def add(p, q, r):  
    temp = p + q + r  
    return temp
```

```
sum = add(1, 10, 100)
```

NOTE: variables are not declared explicitly, they are declared implicitly the first time they are assigned.

E.g. in this example, “temp” is bound to an object that results from evaluating a sum of objects that are defined as integers here. So the function returns an object that is an integer.

```
a, b, c = 1, 10, 100
```

Global variables

```
def printsum():
    print(a + b + c)
```

No arguments
No return value

```
printsum()
```

111

This function does not return any value (no **return** used in the **def**).

You can define variables outside of a function as **global variables**, then refer them inside the function. This is how outer and inner scope works in Python.

Lists

a is a variable (so an object), and this object is bound to a list of 3 integers

```
a = [1, 2, 3]
```

```
print( a[0] )
```

1

Index

```
b = [1, True, a, 'four']
```

```
print(b)
```

[1, True, [1, 2, 3], 'four']

Mixed types

Nested

Slice of a list

```
a = [1, 2, 3]  
print( a[0] )
```

1

Index

```
b = [1, True, a, 'four']  
print(b)
```

[1, True, [1, 2, 3], 'four']

Mixed types

Nested

```
print( b[0:2] )
```

[1, True]

Slice

[Start : Stop : Step]

Function + List

A function that returns a list

```
def sum_and_diff(a, b):  
    return [a + b, a - b]
```

```
x = sum_and_diff(7, 4)
```

That list is an object

```
print(type(x))  
print(x)  
print(x[0])
```

```
<class 'list'>  
[11, 3]  
11
```

Dictionary

```
ages = {'Tom': 32, 'Dick': 58, 'Harry': 28}
```

```
for k in ages.keys():
    print(k, ages[k])
```

Tom 32
Dick 58
Harry 28

```
del ages['Tom']
```

```
ages['Bert'] = 49
```

```
for k,v in ages.items():
    print(k, v)
```

Dick 58
Harry 28
Bert 49

Modules

Note: the 2 modules contains each a function which has the same name! But these are disambiguated by calling them with the module name, plus a “.” and then the function name

`spam.py`

```
def fn():
    print('spam')
```

`eggs.py`

```
def fn():
    print('eggs')
```

```
import spam
import eggs
```

```
spam.fn()
eggs.fn()
```

spam
eggs



Modules

A couple of tricks..

```
import spam  
import eggs
```

```
spam.fn()  
eggs.fn()
```

```
spam  
eggs
```

```
from spam import fn  
from eggs import fn as fn2
```

```
fn()  
fn2()
```

```
spam  
eggs
```

Packages

Packages are simply fancy modules.

```
top/
    __init__.py
    mod1.py
    mod2.py
    part1/
        __init__.py
        mod3.py
        mod4.py
    part2/
        __init__.py
        mod5.py
        mod6.py
```

```
import top.mod1
import top.mod2
import top.part1.mod3
import top.part1.mod4
import top.part2.mod5
import top.part2.mod6

top.part2.mod6.fn()
```

If you want to avoid all this, just do:

```
from top.part2.mod6 import fn
fn()
```

Another example

Do you understand - or at least “feel” - what these few lines may do?

```
from keras.models import Sequential
from keras.layers import Dense, Activation

model = Sequential()

model.add(Dense(units=64, input_dim=100))
model.add(Activation('relu'))
model.add(Dense(units=10))
model.add(Activation('softmax'))

model.compile(loss='categorical_crossentropy',
              optimizer='sgd',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5, batch_size=32)

print(model.evaluate(x_test, y_test, batch_size=128)[1])
```

Classes

List

(this is real - despite basic - code that aims at creating a machine learning model..)

Context manager

The **with** construct. The context manager does not do anything :) it just sets up the following code..

```
with open('spam.txt') as file:  
    for line in file:  
        print(line, end='')
```

That's all for now on Python. What next?

About Python, in general, if you are a beginner:

- install python on your (local) resource (e.g. PC/laptop)
- browse online and pick up a good tutorial
- familiarise with the basic concepts and the syntax
- **exercise, exercise, exercise!** This will be most probably (one of) your “language(s)” at work, so get used to it, and constantly improve!

About Python, in the context this course:

- no installation, focus on understanding the syntax, as we will (sort of) give it for granted
- if you are unsure on what some portion of python code does, set aside, check/study it carefully, then come back to main stream of the course..
 - ❖ you are students at different levels, so this part is left to you - ask if you need assistance!

Applied Machine Learning - Basic

Prof. Daniele Bonacorsi

Hands-on:
Jupyter/Colab

Data Science and Computation PhD + Master in Bioinformatics
University of Bologna



Hands-on: Get familiar with Jupyter

Jupyter notebooks



<https://jupyter.org/>

An open-source web application that allows to create and share code, plots, documents.

It offers **one single environment** for:

- code
- comments on the code
- data analysis + data visualisation
- any additional context (e.g. text, formulas, even media files..)

Perfect for **streamlining an entire workflow**.

And excellent **in the prototyping phase**.

Jupyter notebooks



Other features:

- **pedagogical tasks**: produce a report, a course, a book, anything
- **sharing tasks**: individual analysis but also good for sharing with co-workers

In a nutshell, a powerful, versatile, shareable tool

- not a surprise that is very popular among data scientists

User interface

Any context is organised in independent **cells**.

A cell is executed individually

- a user can test a portion of her/his code without needing to run all code
- naturally fits interpreted languages

Getting started (local resource)

Install

- e.g. anaconda gives you python + jupyter (and plenty more..) in one shot
- pip install jupyter

Launch

- \$ jupyter notebook &
- Jupyter notebook opens up in your default web browser with the URL
 - ❖ <http://localhost:8888/tree>



Your console

List all the files in the dir
you launched “jupyter notebook” from

The screenshot shows the Jupyter Notebook web interface. At the top, there is a navigation bar with the Jupyter logo on the left and a 'Logout' button on the right. Below the navigation bar, there are three tabs: 'Files' (which is highlighted with a black border), 'Running', and 'Clusters'. A black arrow points from the text above to the 'Files' tab. Below the tabs, there is a message 'Select items to perform actions on them.' To the right of this message are three buttons: 'Upload', 'New ▾', and a refresh icon. The main area is a file browser with a sidebar on the left containing checkboxes, a dropdown menu, and a home icon. The list of files and folders is as follows:

- [Anaconda-Navigator.app](#)
- [BACKUP](#)
- [bin](#)

Your console

List all running processes, i.e. terminals and notebooks you currently have open

Provided by iPython parallel

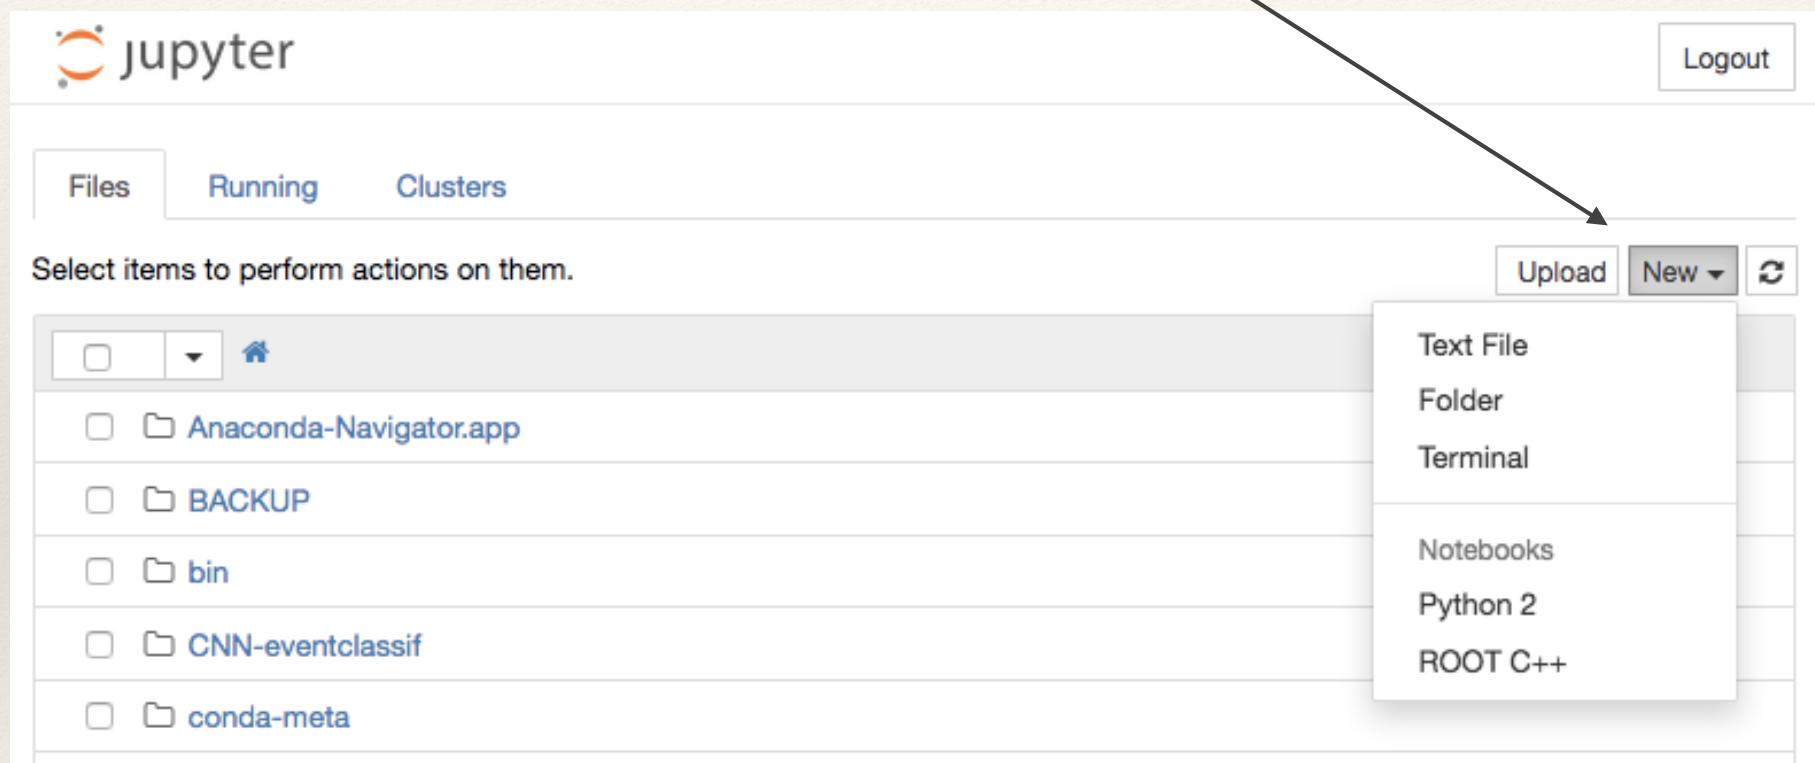
The screenshot shows the Jupyter interface with the following elements:

- Logo:** A Jupyter logo icon followed by the word "jupyter".
- Navigation:** A horizontal menu bar with three items: "Files", "Running" (which is highlighted), and "Clusters".
- User Options:** A "Logout" button in the top right corner.
- Section Header:** "Currently running Jupyter processes" with a refresh icon.
- Terminals:** A section titled "Terminals" with a dropdown arrow. It displays the message: "There are no terminals running."
- Notebooks:** A section titled "Notebooks" with a dropdown arrow. It displays the message: "There are no notebooks running."

Two arrows point from the text "List all running processes, i.e. terminals and notebooks you currently have open" to the "Running" tab and the "Clusters" tab in the navigation bar.

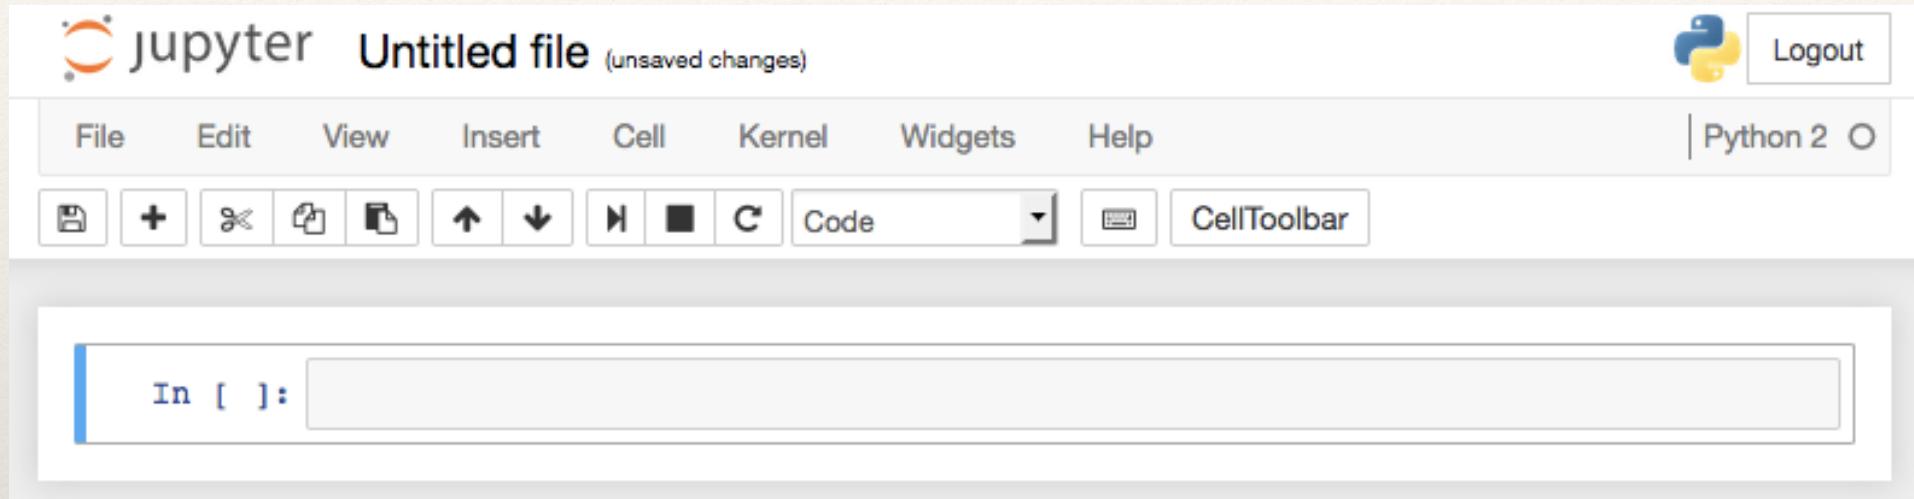
Start using a notebook

Open a new file: Text file, Folder, Terminal, plus **notebooks** in various languages



E.g. in this example, you could select “Python 2” and open a notebook where code cells are interpreted if you write code in Python 2 syntax into them

As simple as this



We will take some time later on, and you can play around it yourself!

Basic options

Explore it yourself. Documentation is good. Focus on two things.

Dropdown menu

- **Code**: self-explanatory; it is where you type your code
- **Markdown**: where you type your text, e.g. comments, notes, ..
- **Raw NBConvert**: a command line tool to convert your notebook into another format (like HTML)
- **Header**: you can use text and do the same.. in most recent versions this will be disappearing and shading into Markdown itself

Command palette

- shortcut commands: learning a few will suffice.

How many keyboard shortcuts do I need?

Few people know many of them.. most people knew just few

One can survive with B/A, DD, <ESC>+M/Y:

- **B/A:** position yourself and add a cell below/above the current one (B or A respectively)
- **DD:** delete a cell
- **<ESC>+M/Y:** change mode from code (Y) to markdown (M), and viceversa

plus a minimal adoption of easy and self-explanatory GUI items

- e.g. File → Open → ..

Not only Python

We showed an example interface allowing for Python (2).

Jupyter is not limited to Python. You can use other languages in Jupyter notebooks

- e.g. R, Julia, JavaScript, etc
 - ❖ suppose you like the 'ggplot2' package in R: you can use it within a Jupyter notebook!

You need to enable the relevant Jupyter "kernel"

- e.g. to enable R in Jupyter, you need the 'IRKernel' (a dedicated kernel for R) which is available on github.
- e.g. for Javascript, you need the 'IJavascript' kernel
- same for other languages..

More [1/2]: Magic functions

Many exist, do explore..

- they can be used line-wise and cell-wise

```
In [3]: %lsmagic
Out[3]: Available line magics:
alias %alias_magic %autocall %automagic %autosave %bookmark %cat
%cd %clear %colors %config %connect_info %cp %debug %dhist %dirs
%doctest_mode %ed %edit %env %gui %hist %history %killbgscripts %
ldir %less %lf %lk %ll %load %load_ext %loadpy %logoff %logon %
logstart %logstate %logstop %ls %lsmagic %lx %macro %magic %man
%matplotlib %mkdir %more %mv %notebook %page %pastebin %pdb %pdef
%pdoc %pfile %pinfo %pinfo2 %popd %pprint %precision %profile %pr
un %psearch %psource %pushd %pwd %pycat %pylab %qtconsole %quickr
ef %recall %rehashx %reload_ext %rep %rerun %reset %reset_selectiv
e %rm %rmdir %run %save %sc %set_env %store %sx %system %tb %
ime %timeit %unalias %unload_ext %who %who_ls %whos %xdel %xmode

Available cell magics:
%%! %%HTML %%SVG %%bash %%capture %%debug %%file %%html %%javascr
ipt %%js %%latex %%perl %%prun %%pypy %%python %%python2 %%python
3 %%ruby %%script %%sh %%svg %%sx %%system %%time %%timeit %%wri
tefile

Automagic is ON, % prefix IS NOT needed for line magics.
```

- play around e.g. with %time and %%time

```
In [16]: %time sum(range(100))
```

```
In [19]: %%time
sum(range(100))
sum(range(100000))
```

More [2/2]: Interactivity via ipywidgets

Add more interactivity - sometimes an added value:

```
from ipywidgets import widgets
```

Many examples online - explore!



Hands-on:
Get familiar with Google Colab

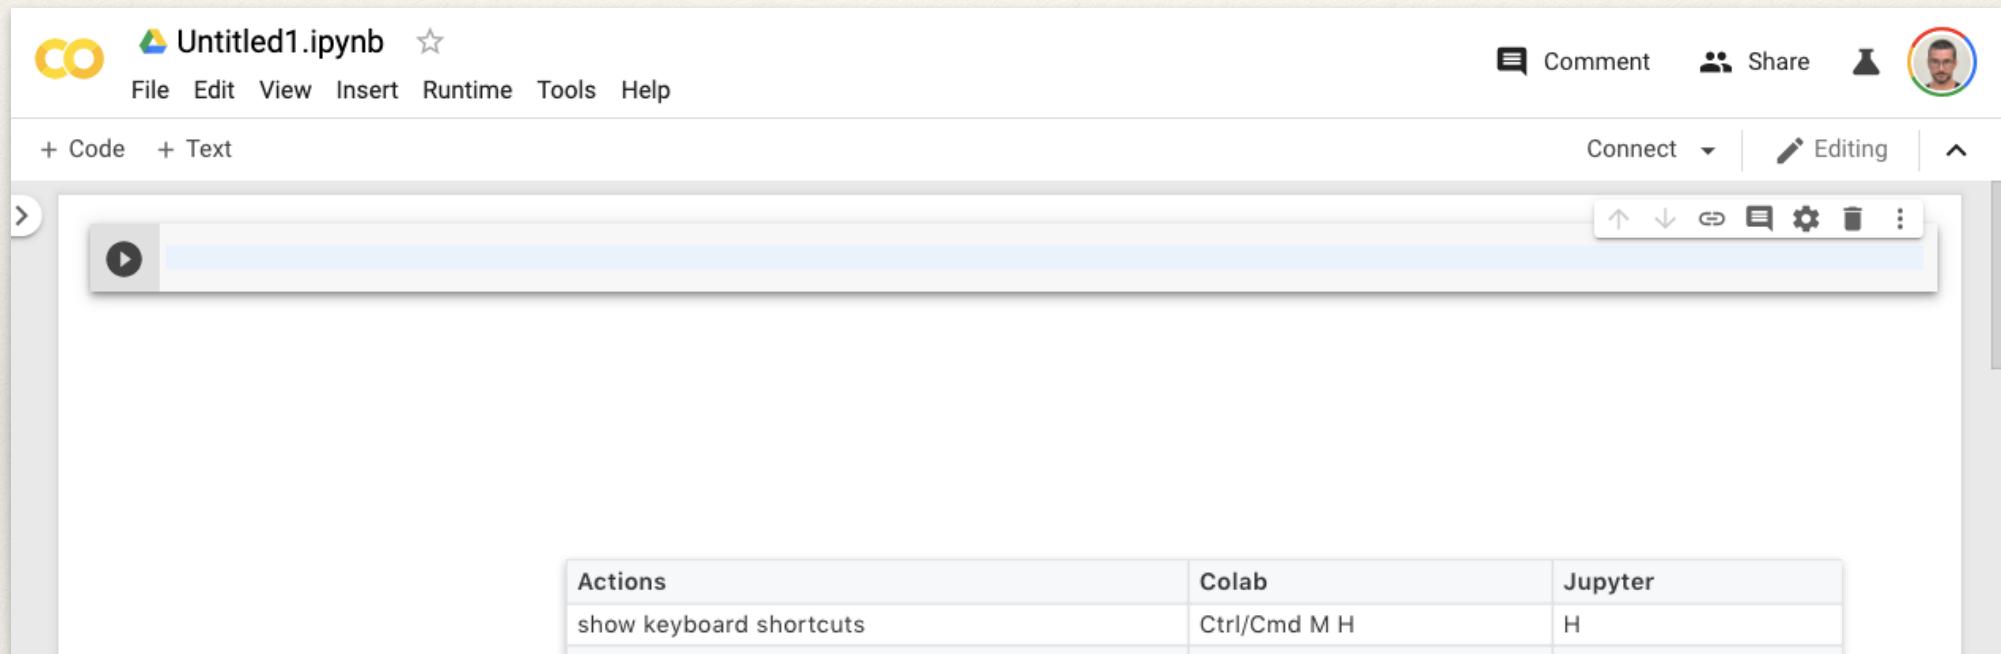
Google Colab(oratory)

colab

<https://colab.research.google.com/>

(allow me oversimplifications here..)

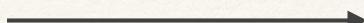
In a nutshell: **Jupyter notebooks on the cloud.**



The screenshot shows the Google Colab interface. At the top, there's a navigation bar with a 'CO' logo, the file name 'Untitled1.ipynb', and icons for Comment, Share, and profile. Below the bar are buttons for '+ Code' and '+ Text'. The main area contains a code cell with a play button icon. To the right of the cell are controls for moving up/down, running, and deleting. A large table below compares keyboard shortcuts for Colab and Jupyter across several actions.

Actions	Colab	Jupyter
show keyboard shortcuts	Ctrl/Cmd M H	H
Insert code cell above	Ctrl/Cmd M A	A
Insert code cell below	Ctrl/Cmd M B	B
Delete cell/selection	Ctrl/Cmd M D	DD
Interrupt execution	Ctrl/Cmd M I	II
Convert to code cell	Ctrl/Cmd M Y	Y
Convert to text cell	Ctrl/Cmd M M	M
Split at cursor	Ctrl/Cmd M -	Ctrl Shift -

Keyboard shortcuts:
Google Colab vs Jupyter



How to get started on Colab

colab

Quickest path to get started:

- **create** a new Python notebook, no installation, just **code**
 - ❖ import modules, code what you want, execute cells, ..
 - ❖ IMPORTANT: most used data sciences stuff is already installed ("no pip, just import")

Proper way to get started: **remember you are on the cloud!**

- **create** a new Python notebook
 - ❖ or **open an existing notebook** from gdrive, github, etc.. or upload your ipynb from your laptop
- **Save** a copy on your own gdrive, and **code** from that one
 - ❖ all future changes will be saved

Best of “Jupyter + Google drive”?

colab

Google Colab’s GUI is (intentionally) very similar to Jupyter’s, plus features typical of Google Drive:

- you can share and use the notebooks like regular Google Docs
- you can create handy widgets using special comments in your code, and more

Colab: select your Runtime

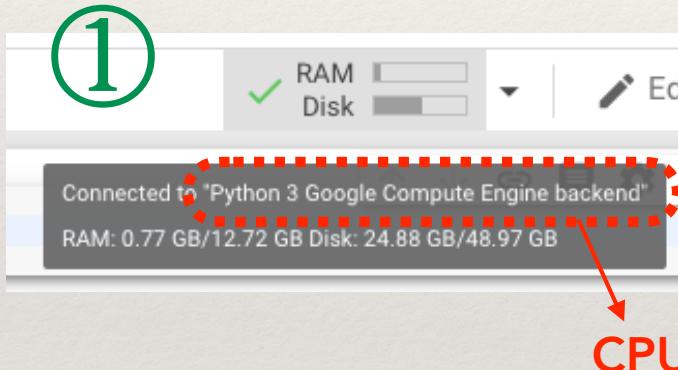
colab

Once you open a Colab notebook,
you get a Colab **Runtime**

- this is the free Google VM
dedicated to you for a while

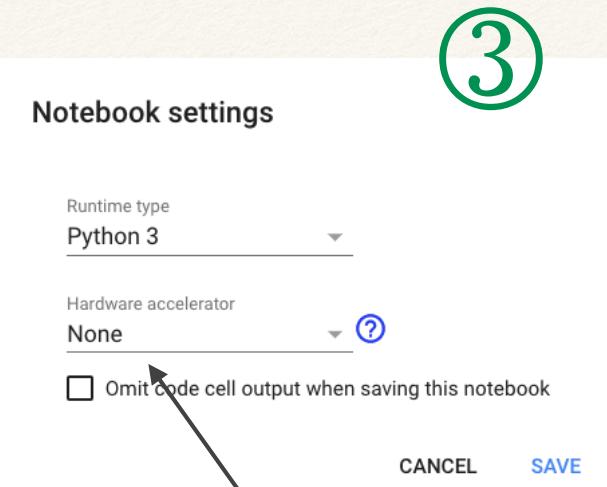
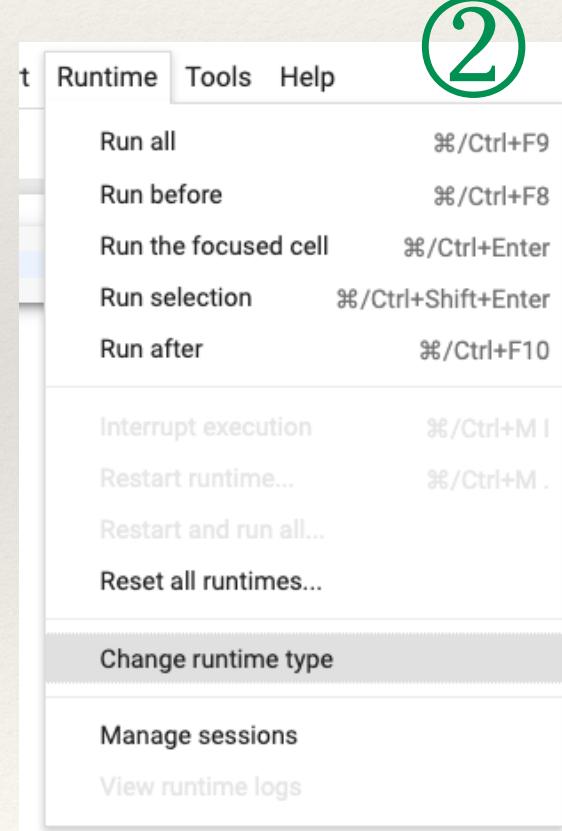
By default, the initial Runtime is CPU-only

- but you can change this:



colab VM states are:

Connecting,
Allocating,
Initialising, and its
actual RAM/Disk
monitoring widget



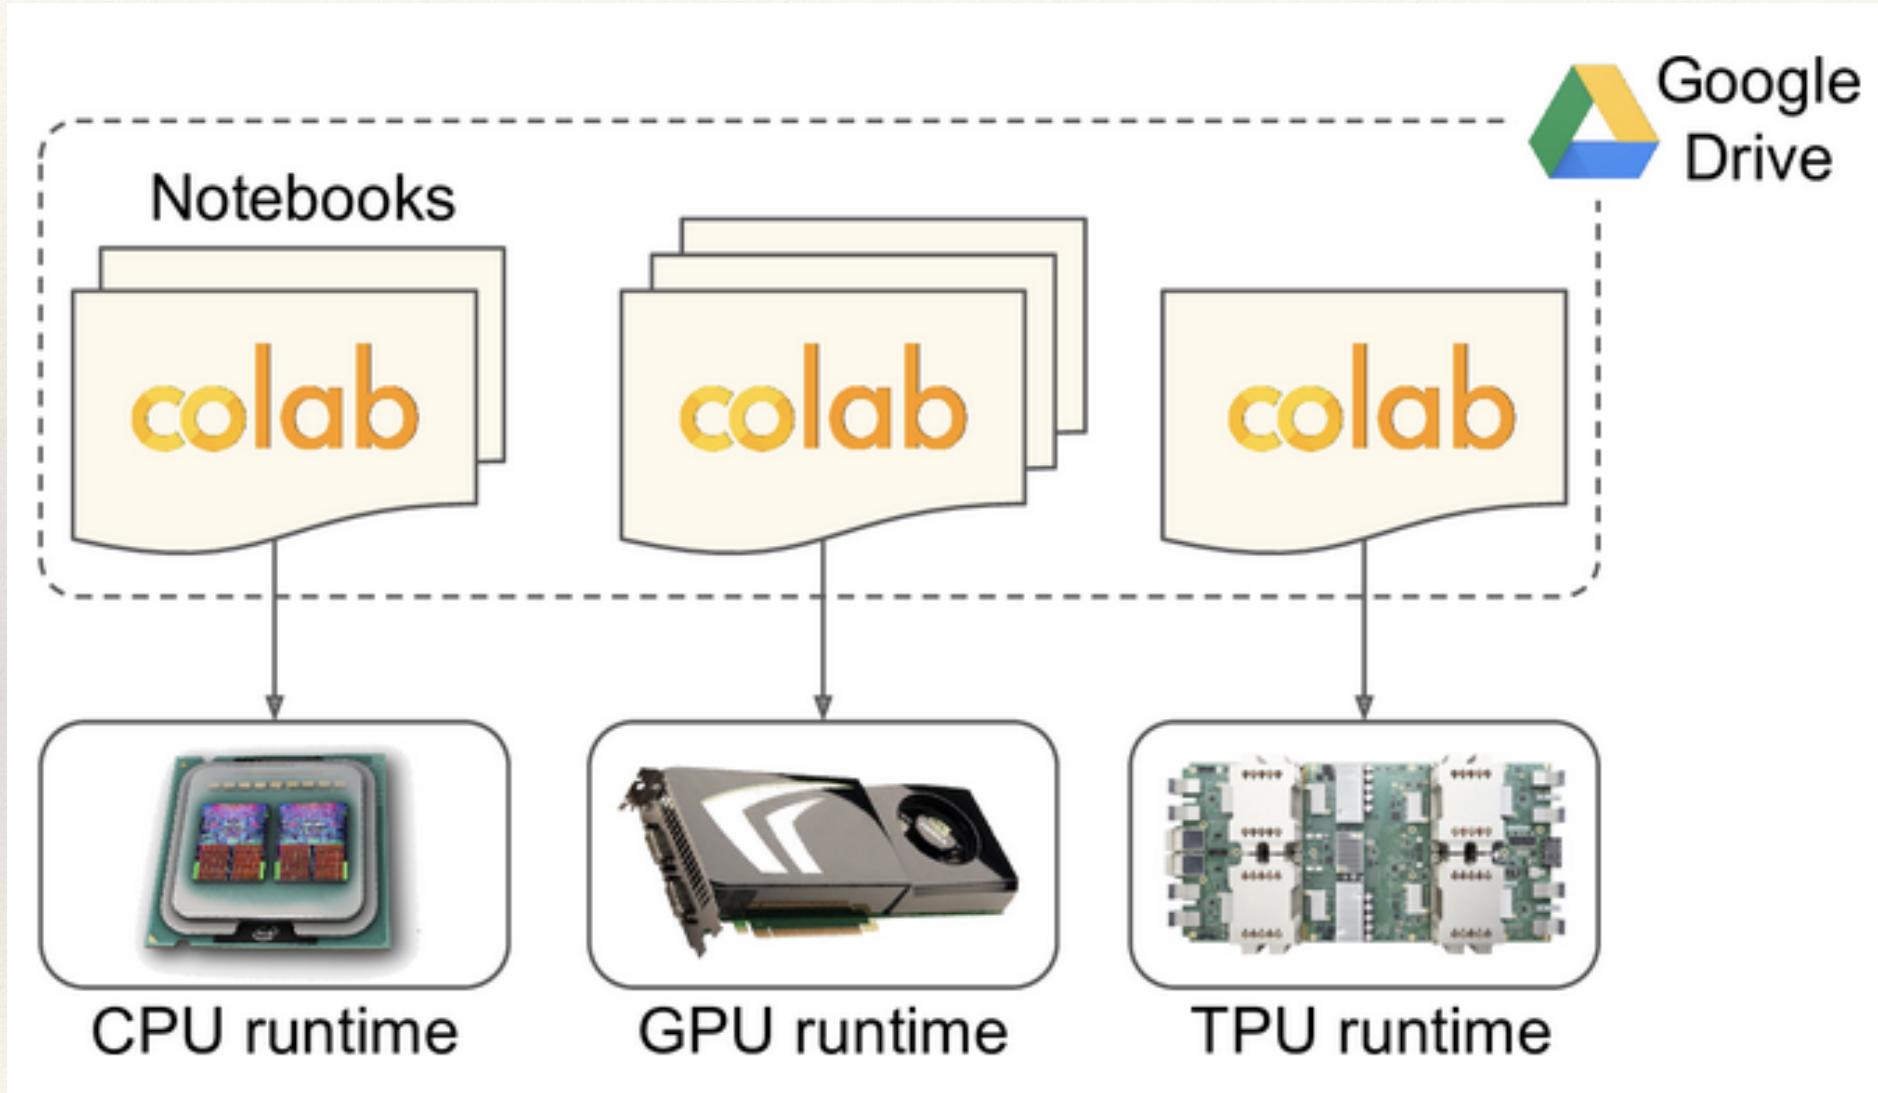
select GPU here



(NOTE: yes, also TPUs are available..)

Colab Runtimes

colab



Lifetime of a Colab Runtime



As from the Colab FAQs:

- “Colaboratory is intended for interactive use. Long-running background computations, particularly on GPUs, may be stopped. Please do not use Colaboratory for cryptocurrency mining.”

The web interface will automatically disconnect you from the Colab Runtime after **~30 minutes** of unattended connection

- you can have it again, but when you reconnect to the Colab Runtime, it may have been reset (so make sure you download what you needed before moving to something else)

Even if you stay connected and actually code, the Colab Runtime will automatically shut down after **12 hours**

- this is meant to prevent long-running computations

Still a neat tool for experimenting on GPUs (and for free)..

Which GPU am I getting in Colab?

Tesla K80, 12 GB RAM,
up to 12 hrs in a row.



[DISCLAIMER: this information might be outdated in the future, pretty soon..]

Which GPU on Colab



Actually:

- the free version of **Colab** mostly provides Tesla K80 GPU
- **Colab pro**, the paid version (\$9.99/month (*)), provides access to Tesla T4 or P100 GPUs and also other benefits like longer runtimes and priority access to TPUs

(* as of March 2023

Additional material

Not only GPU..
the Google TPU

Specialised hw implementations for Deep Learning

— Google's (proprietary) TPU —

2006: Google engineers started working on deploying GPUs, FPGAs, or custom ASICs in their Data Centres (only few applications could effectively run on them, though)

2013: projection that if 100% people used voice search for 3 mins/day using speech recognition DNNs, it would have required Google's data centres to double in size in order to meeting computation demands

Google started a project whose mandate was improve cost-performance by 10x over GPUs

- quickly produce a custom ASIC for inference
- bought off-the-shelf GPUs for training

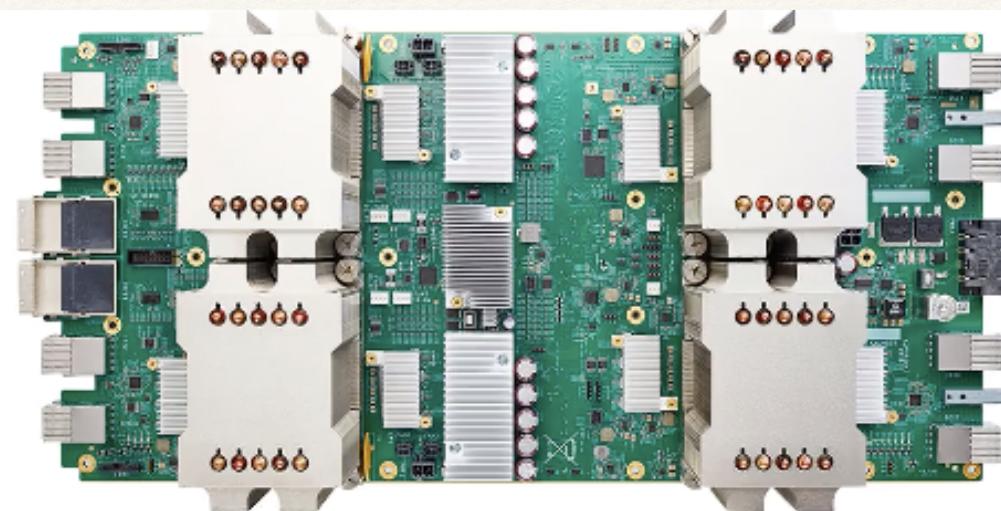
The **Tensor Processing Unit (TPU)** - Google's first custom ASIC to be used in the domain of fast inference of DNNs, programmed using the Tensorflow framework - was designed, verified, built, and deployed (first in 2015) in just 15 months (!)

Google's TPU v1 vs v2



TPU v1

Launched in 2015
Inference only



TPU v2

Launched in 2017
Inference and training

 Search Search ranking Speech recognition	 Translate Text, graphic and speech translation	 Photos Photos search	
---	---	--	---

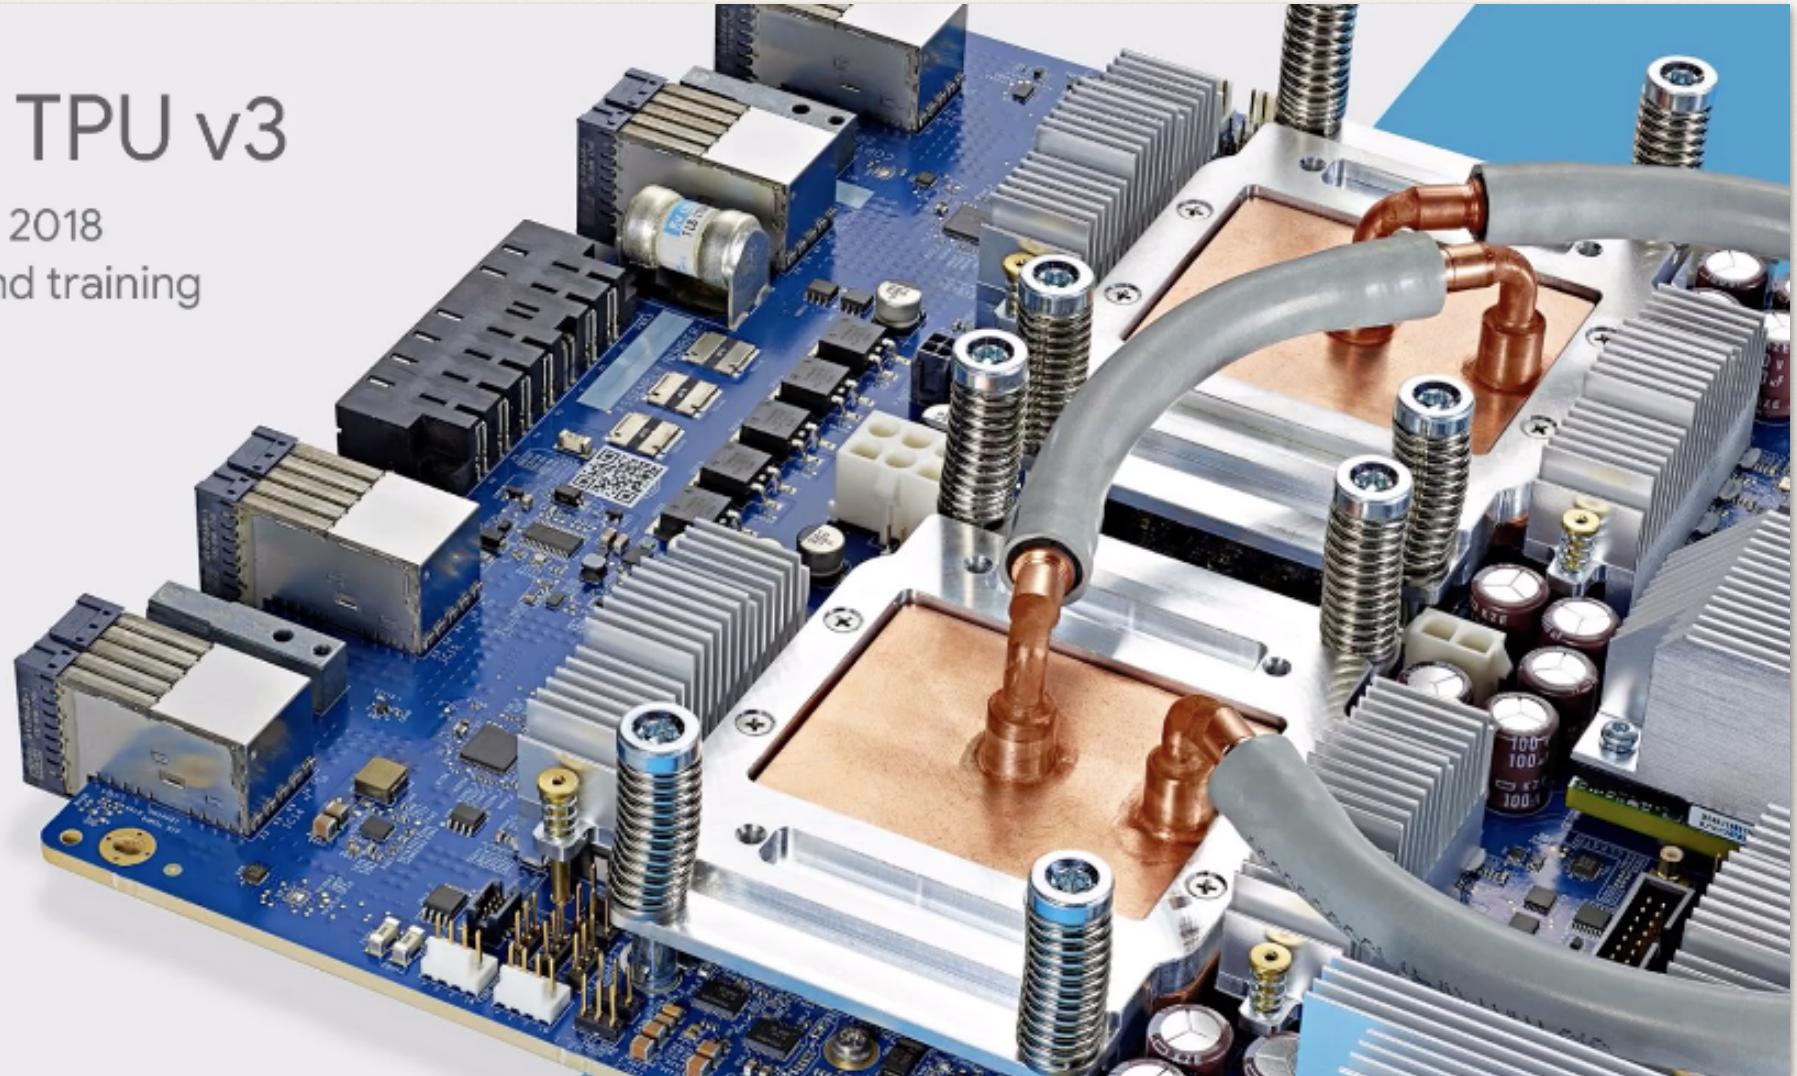
v2: both internally
and on Cloud

Google's Cloud TPU v3

Cloud TPU v3

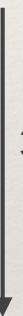
Launched in 2018

Inference and training



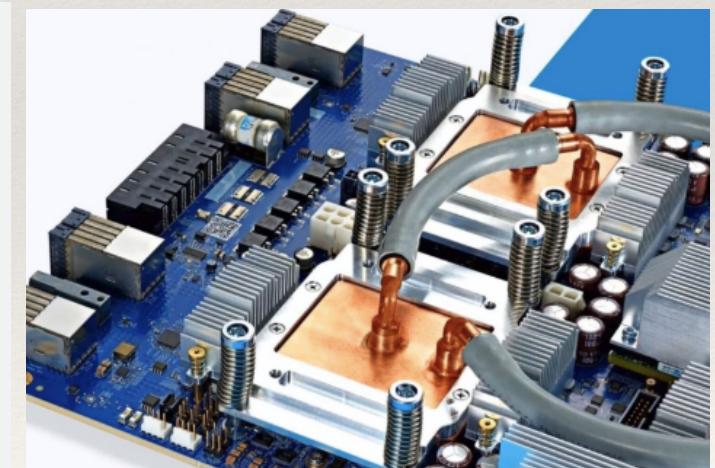
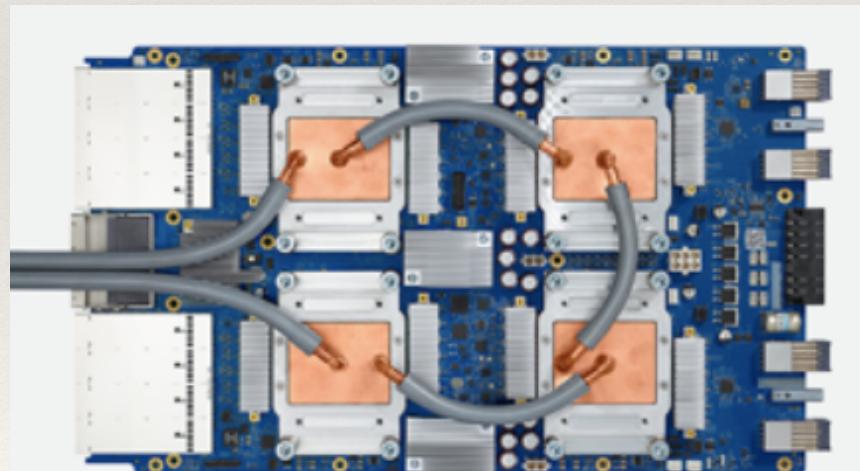
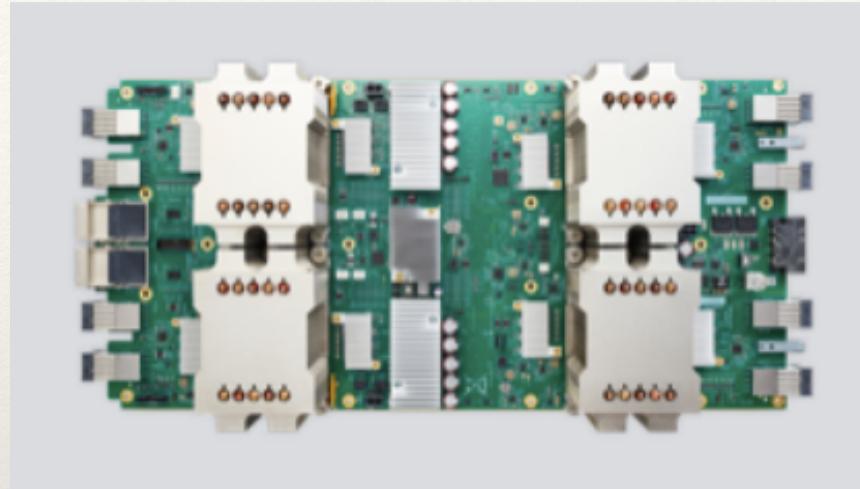
Google's TPU v2 vs v3

Cloud TPU v2
180 teraFLOPS
64 GB High
Bandwidth Memory



>2x teraFLOPS

Cloud TPU v3
420 teraFLOPS
128 GB HBM



Interconnect TPUs...

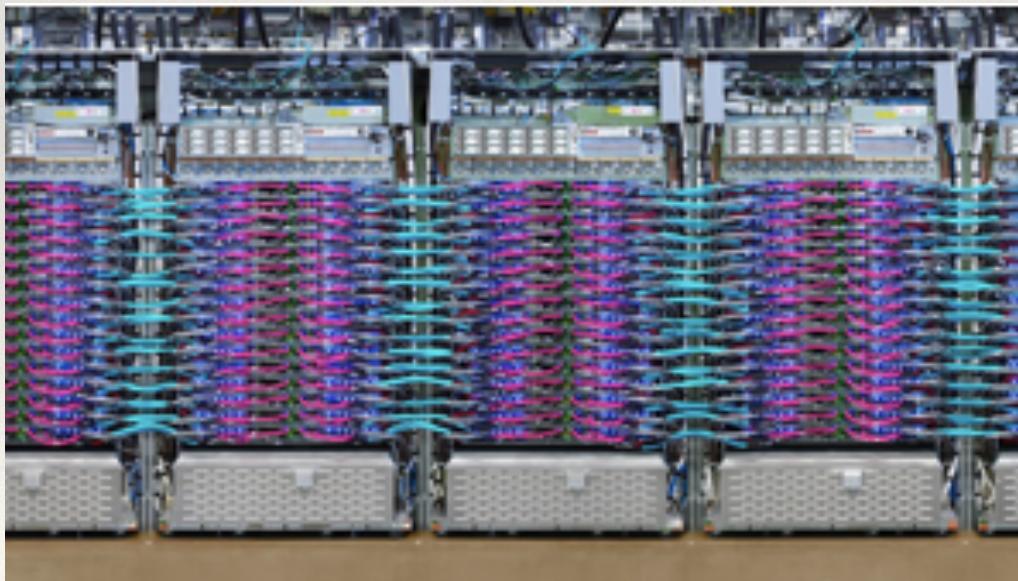


Cloud TPU v2 Pod (beta)

11.5 petaflops

4 TB HBM

2-D toroidal mesh network



Cloud TPU v3 Pod (beta)

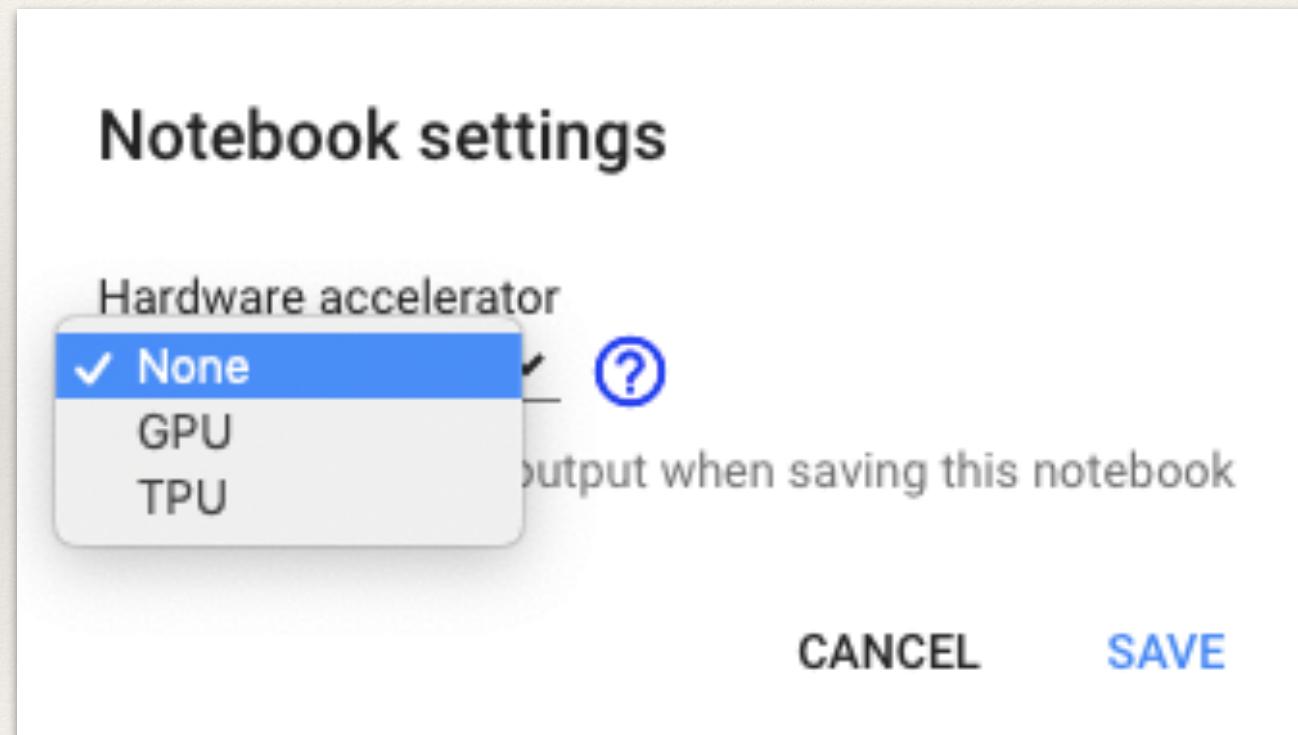
100+ petaflops

32 TB HBM

2-D toroidal mesh network

You can actually choose:

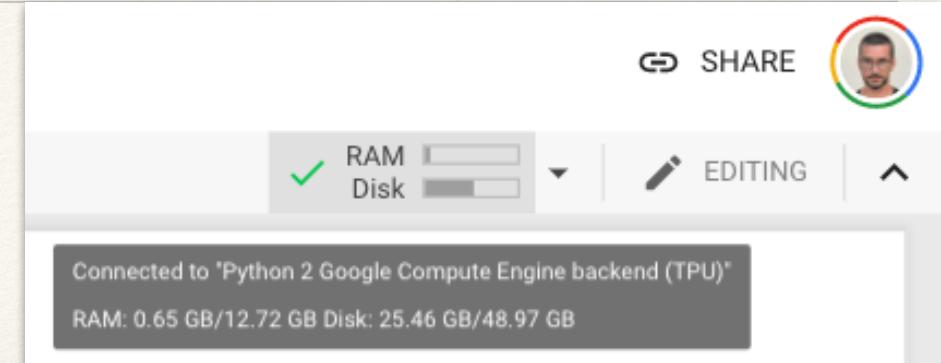
- no accelerator
- GPU
- TPU



"Why am I seeing TPU slower than GPU?"

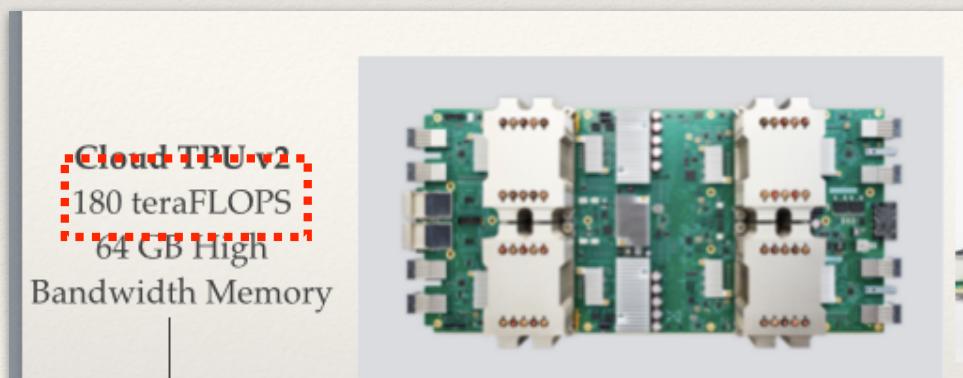
Testing != Benchmarking/Profiling ..

be careful about your conclusions!



Try some floating point computations that give you the (T)FLOPS

- e.g. <https://colab.research.google.com/notebooks/tpu.ipynb>



```
[3] N = 4096
COUNT = 100
import time

def flops():
    x = tf.random_uniform([N, N])
    y = tf.random_uniform([N, N])
    def _matmul(x, y):
        return tf.tensordot(x, y, axes=[[1], [0]]), y

    return tf.reduce_sum(
        tf.contrib.tpu.repeat(COUNT, _matmul, [x, y]))
)

tpu_ops = tf.contrib.tpu.batch_parallel(flops, [], num_shards=8)

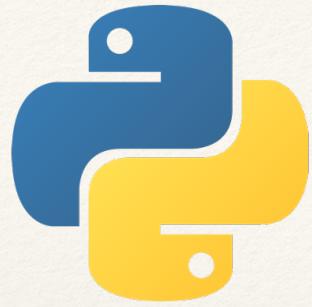
session = tf.Session(tpu_address)
try:
    print('Warming up...')
    session.run(tf.contrib.tpu.initialize_system())
    session.run(tpu_ops)
    print('Profiling')
    start = time.time()
    session.run(tpu_ops)
    end = time.time()
    elapsed = end - start
    print(elapsed, 'TFlops: {:.2f}'.format(1e-12 * 8 * COUNT * 2*N*N*N / elapsed))
finally:
    session.run(tf.contrib.tpu.shutdown_system())
    session.close()

Warming up...
Profiling
0.6749346256256104 TFlops: 162.91
```

If you want to know more about a TPU

Google TPU demo

[URL: <https://storage.googleapis.com/nexttpu/index.html>]



Hands-on: Python 101 notebook on Jupyter

Name ↑	Drive
 AML2223Bas_Lectures	
 AML2223Bas_Notebooks	
 AML2223Bas_Lectures 	
 AML2223Bas_StudentsDirectory 	

Python101.ipynb

Exercise, and keep this as a
bonus for any future reference!

1110001110100100110011010010100111010100100100101
01110010101001010101010110100101010101011100011
10100100110011010010100111010010010010101110010
101110**Logistic**00**Regression**01001101010010111010
1010010101001101010011001001010111010101001010
10101011010010101010111000111010010011001101010
0101001110101101101011100101001010101010101101
000101010111000111010010011001101001010011101
01001001000**Classification**10 and 10 **representation**
001011101110001110100100110011010010100111010
10010010010**Classification**110111100011101001001
10010001000011111010110100010101011100011101
0010011001100010100000100110011010010100011011
110001001100110100101001110101110010110101010
1011101001101110101010011101011101011010101010

In **classification** problems, the variable **y** that you want to predict tells you if your example belongs to a “class” or not.

We will discuss an algorithm called **logistic regression**

- one of the most popular and most widely used learning algorithms today

Note: more on the its name later..

Classification

Examples:

1. email spam classification: spam / not spam?
 - Example: mail spam classifier, $x(i)$ may be some features of a piece of email, labels are $y \in \{0,1\}$, i.e. y may be 1 if it is spam mail (positive class), and 0 otherwise (negative class)
2. tumors: malignant/benign?
3. classifying online transactions: fraudulent y/n?
4. .. many more! ..

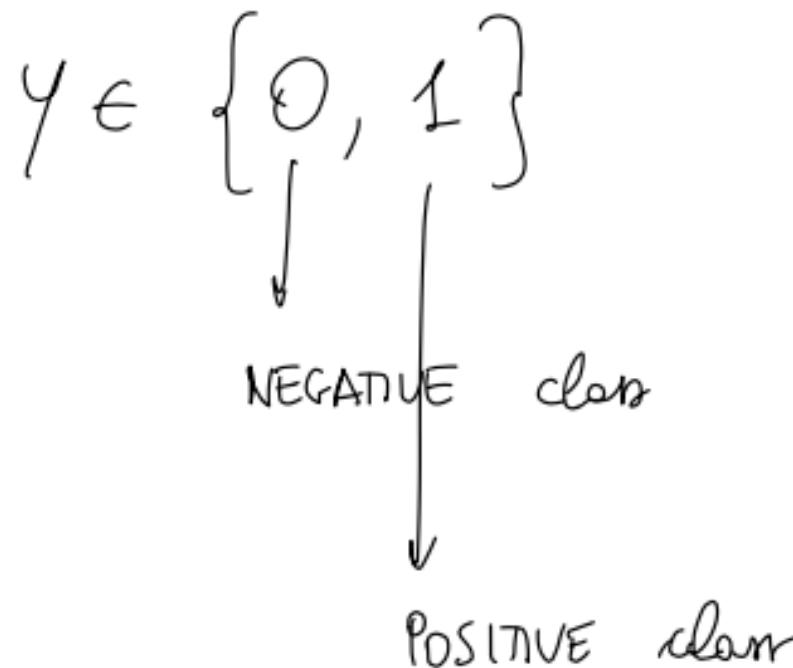
What is common among all these problems?

- the variable that we are trying to predict is a **y** that can take on two values

Basically, **0** or **1**. Two **classes** (i.e. “classification”). A **negative class** and a **positive class**.

Classification

BINARY
CLASSIFICATION



MULTICLASS
CLASSIFICATION

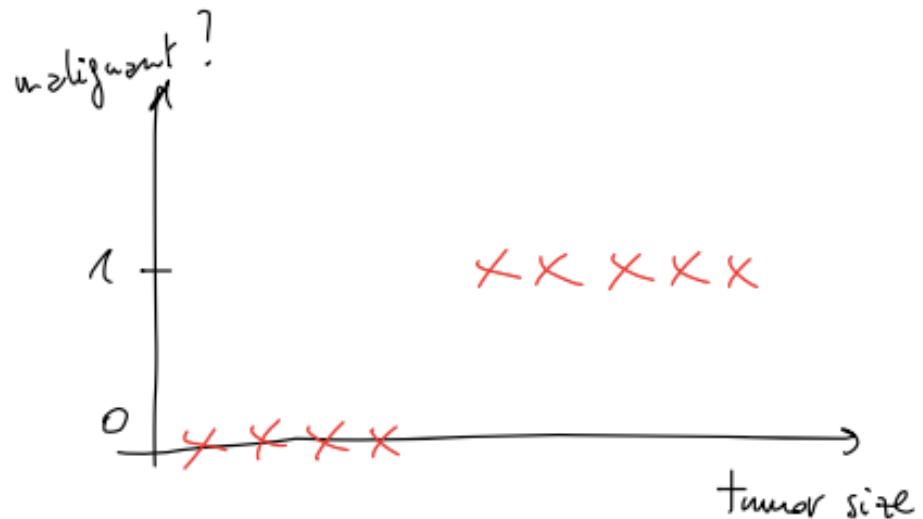
$$y \in \{0, 1, \dots, n\}$$

Classification examples

NOTE: The assignment of the 2 classes - **negative** and **positive** - to **0** and **1** is somewhat arbitrary

- often there is an intuition that a negative class is conveying the absence of something..
 - ❖ (e.g. the absence of a malignant tumour)
- .. whereas the positive class is conveying the presence of something
 - ❖ (e.g. the presence of a malignant tumour - i.e. like in medical exams in general, actually, where you are positive if you indeed have something)
- we will follow some intuition in this, but to be clear: it doesn't matter

Applying linear regression to a classification problem

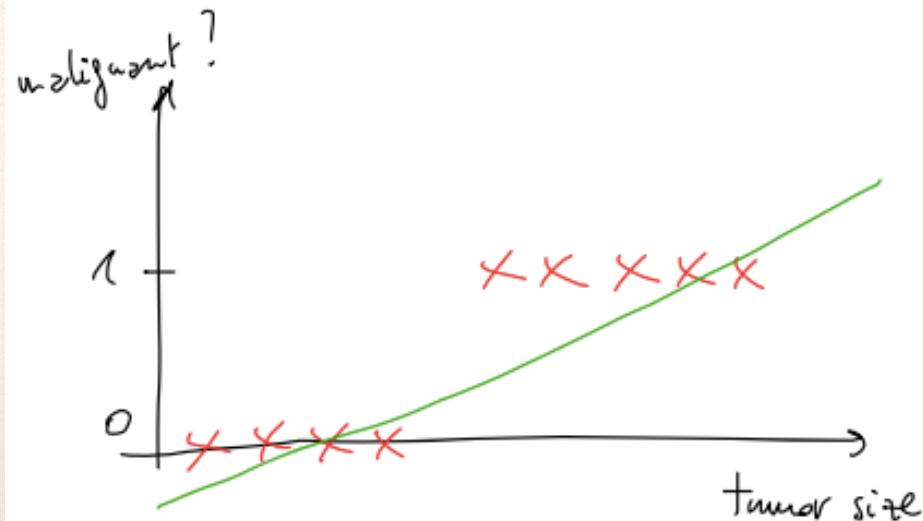


*NOTE: not using different symbols,
as their "true" value is known by
where the data point is on the plot*

Let's try to apply linear regression to this problem.

Do you think it would work?

Applying linear regression to a classification problem

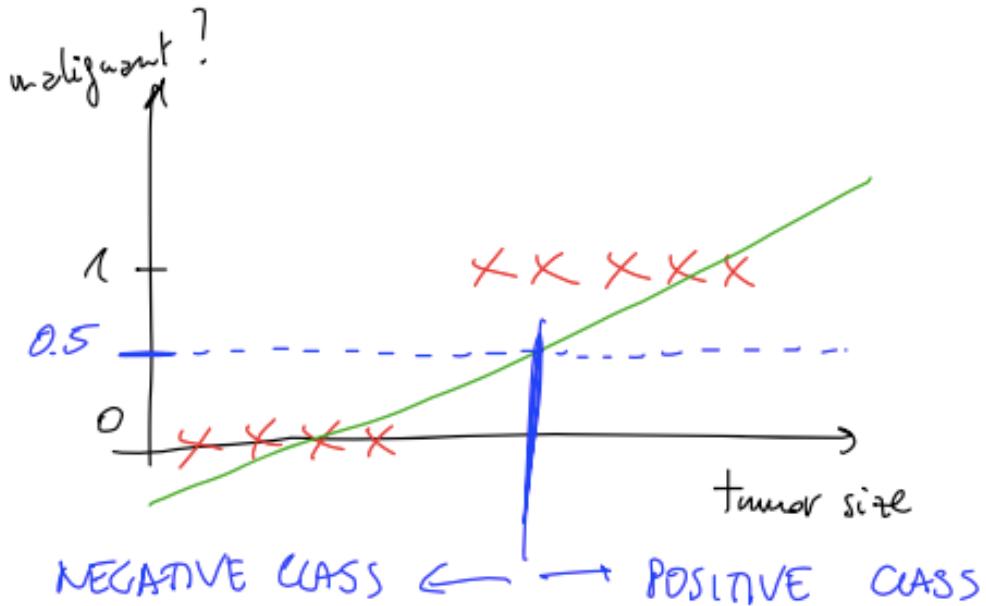


assume you run a linear repr.

$$h_{\theta}(x) = \theta^T x$$

How do you predict?

Applying linear regression to a classification problem



assume you run a linear repr.

$$h_{\theta}(x) = \theta^T x$$

How do you predict?

Threshold classifier output $h_{\theta}(x)$ at 0.5 :

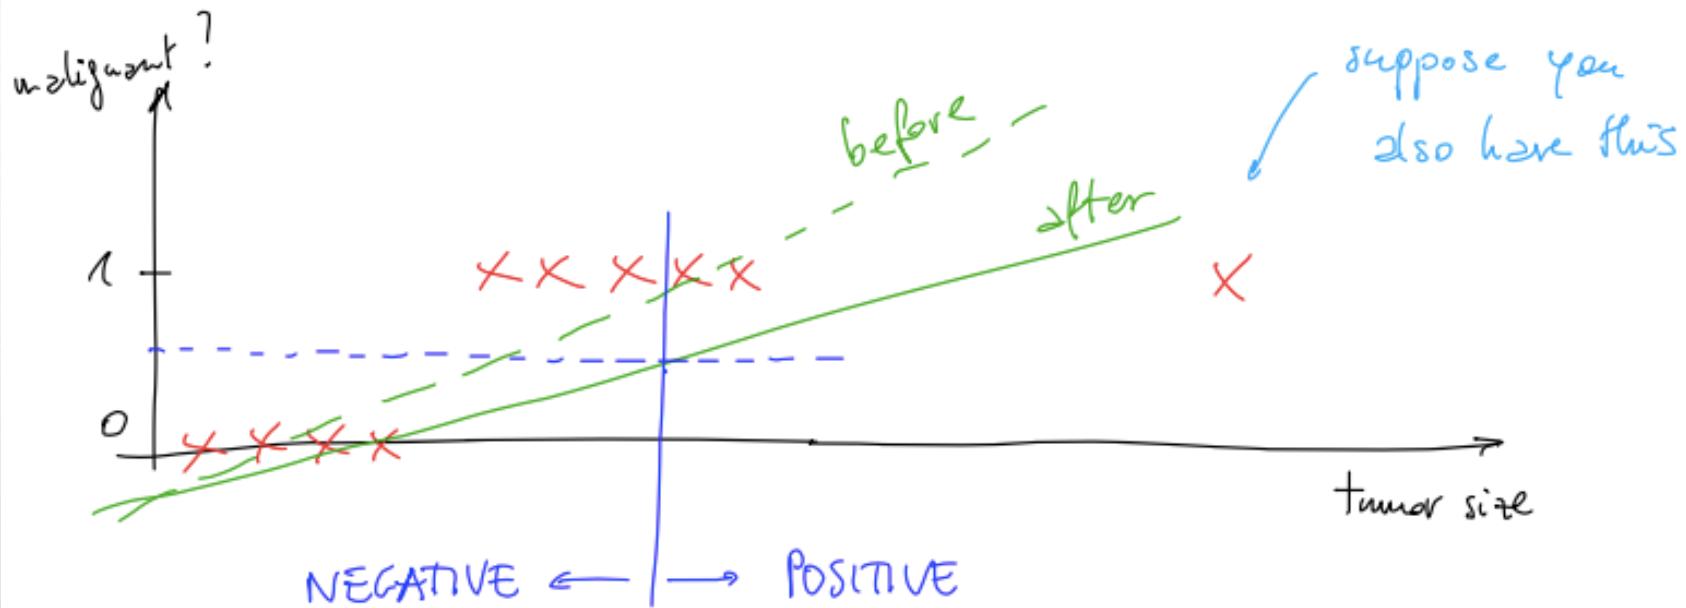
$$h_{\theta}(x) \begin{cases} \geq 0.5 & \Rightarrow \text{predict } "y=1" \\ < 0.5 & \Rightarrow \text{predict } "y=0" \end{cases}$$

Is this good enough? And what if I have more data?

Applying linear regression to a classification problem

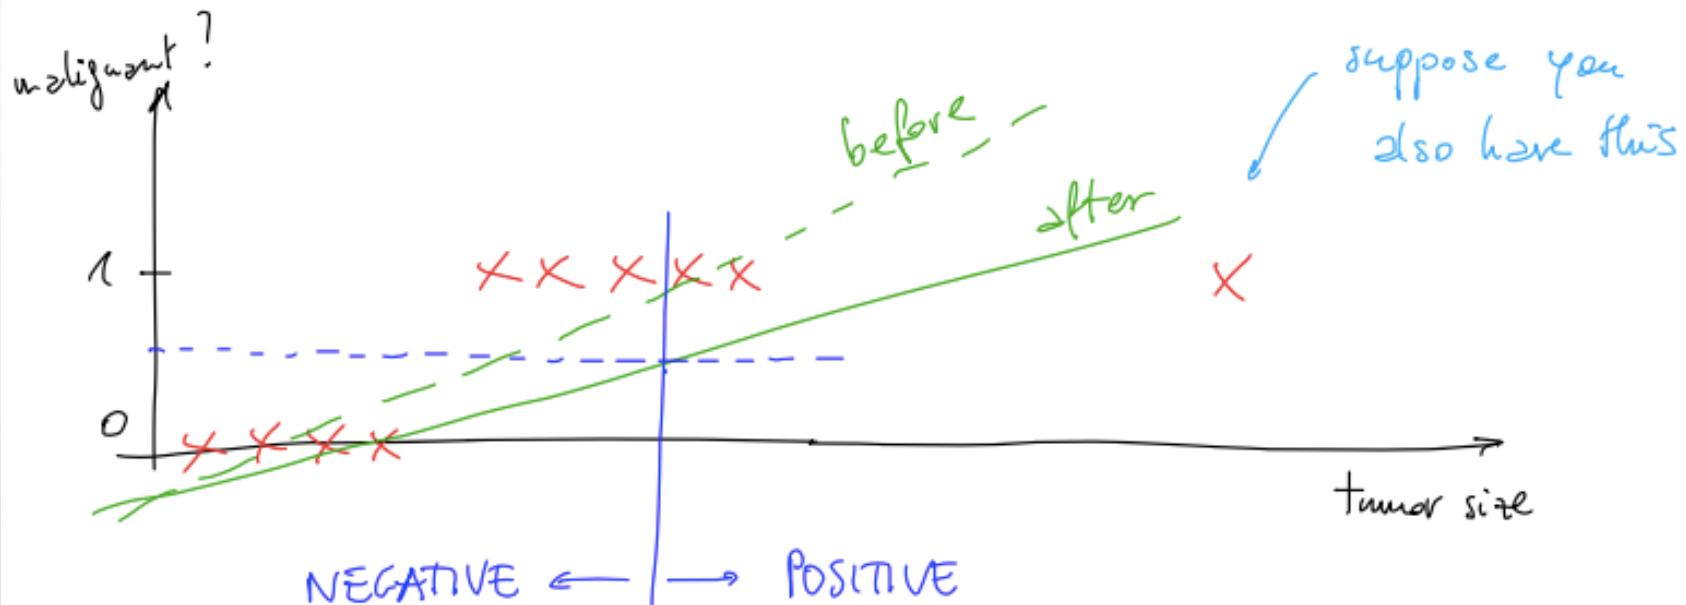


Applying linear regression to a classification problem



One example (that was clear) caused a mess !

Applying linear regression to a classification problem



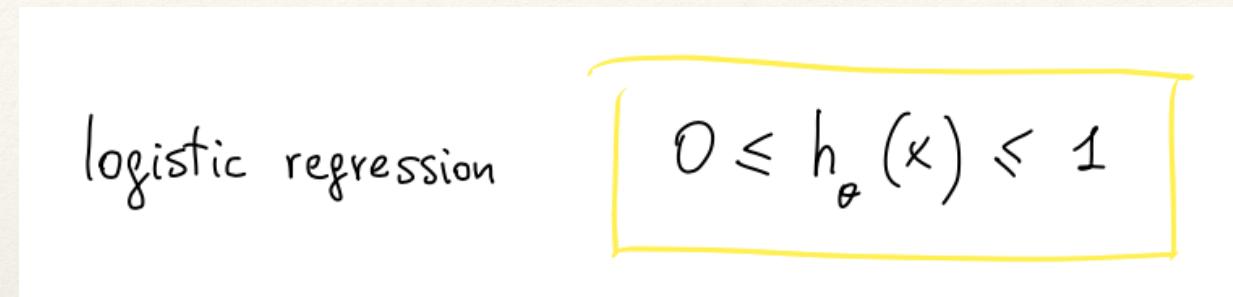
One example (that was clear) caused a mess !

Applying linear regression to a pure classification problem usually is **not** a promising approach

- additionally, your hypothesis can also yield values <0 and >1 , weird for a binary classification problem (should boil down to either 0 or 1, nothing in-between, and for sure nothing outside the interval $[0,1]$!)

Logistic regression

We will introduce a simple algorithm called



And **logistic regression** is - and we will use it as - a **classification** algorithm.

NOTE: It is maybe sometimes confusing that the term “regression” appears in this name even though logistic regression is actually a “classification” algorithm.

- But that's just a name it was given for historical reasons (you'll see). So don't be confused..
- See it like this: a classification problem is just like a regression problem, except that the values you want to predict take on only a small number of discrete values
 - ❖ e.g. predict the probability of raining (regression) but then take the umbrella only if prob > 0.8 (i.e. use the result of regression as a classification for your purposes)

Logistic regression is a classification algorithm that we apply to settings where the label y has discrete values, e.g. when it is either 0 or 1.

Summary

So far, we have a binary classification problem in which y could take on only two values, 0 and 1

- most of what we said will also generalise to the multiple-class case

A first attempt at classification was towards trying linear regression

- i.e. mapping predictions greater than 0.5 as a 1 and smaller than 0.5 as a 0.

However, this method did not work well: classification is actually not dealt with well by a linear function

- As we tried to approach the classification problem just ignoring the fact that y is discrete-valued and using a linear regression algorithm to try to predict y given x as if it was a regression problem, we soon realised how easy it is to construct examples where this approach performs very poorly.
- In addition, we would have been using an approach that allow h to be larger than 1 or smaller than 0, when we know that it could be $y \in \{0, 1\}$

It seems we may benefit indeed from something like **logistic regression** to attack classification problem.

1110001110100100110011010010100111010100100100101
01110010101001010101010110100101010101011100011
10100100110011010010100111010010010010101110010
10**Logistic**00**Regression**010011010100101110101010
0101010011010100110010010101110101010010010101010
1011010010101010111100011101001001100110100101
0011101011011010111001010100101010101011010001
0101011100011101001001100110010100111010100
1001010**Classification**10and10**representation**0010
1110111100011101001001100110100101001110101001
0010010**Hypothesis**10**representation**0011001000100
00111110101101000101010111000111010010011001
1000101000001001100110100101000110111100010011
001101001010011101011100101010101011101001
101110101010011101011101010101010101010110010

Logistic regression: hypothesis representation

I.e. what is the function \mathbf{h} we can use to represent our hypothesis in the case of a logistic regression problem?

In other words, we need to move beyond this:

$$h_{\theta}(x) = \theta_0x_0 + \theta_1x_1 + \dots + \theta_nx_n = \boldsymbol{\theta}^T \mathbf{x}$$

For logistic regression, we will be trying a function \mathbf{g} of \mathbf{h} above..

Logistic regression: sigmoid

LOGISTIC REGRESSION MODEL

I want $0 \leq h_{\theta}(x) \leq 1$

$$h_{\theta}(x) = g(\theta^T x)$$



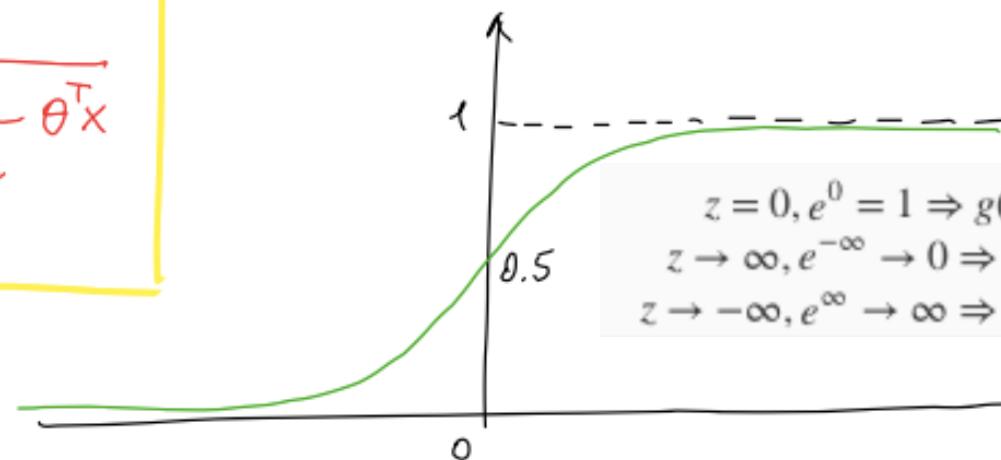
$$g(z) = \frac{1}{1 + e^{-z}}$$

SIGMOID
FUNCTION

(or LOGIST(\mathbf{z}))

origin of the name...

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$



Logistic regression: $g(h)$ is the new h

Moving from the “old” h to “new” $g(h)$ in classification tasks requires some thinking to understand:

How do we interpret now the h ?

- before, h was my guess for y values.. and now?

How to re-interpret the h output for the classification task

$h_{\theta}(x)$ = estimated probability to get $y=1$ on a new input example x

example : Tumors, $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ \text{tumorsize} \end{bmatrix}$ feature vector for a patient

$h_{\theta}(x) = 0.7$ means that patient has 70% prob. that tumor is malignant

You can formalize this as :

$$h_{\theta}(x) = P(y=1 | x; \theta)$$

"my hypothesis can be interpreted as probability that $y=1$, given features x , parameterized by θ "

Deal with both positive/negative classes

We need to state something for 0 too, so:

$$P(y=0 | x; \theta) + P(y=1 | x; \theta) = 1$$

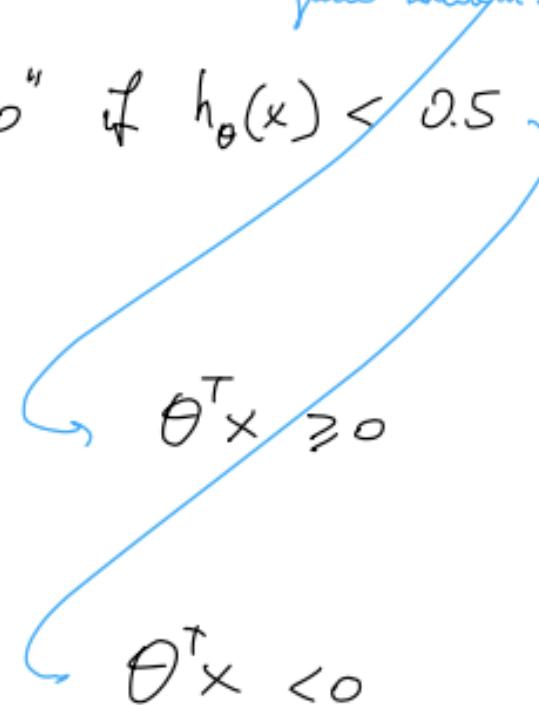
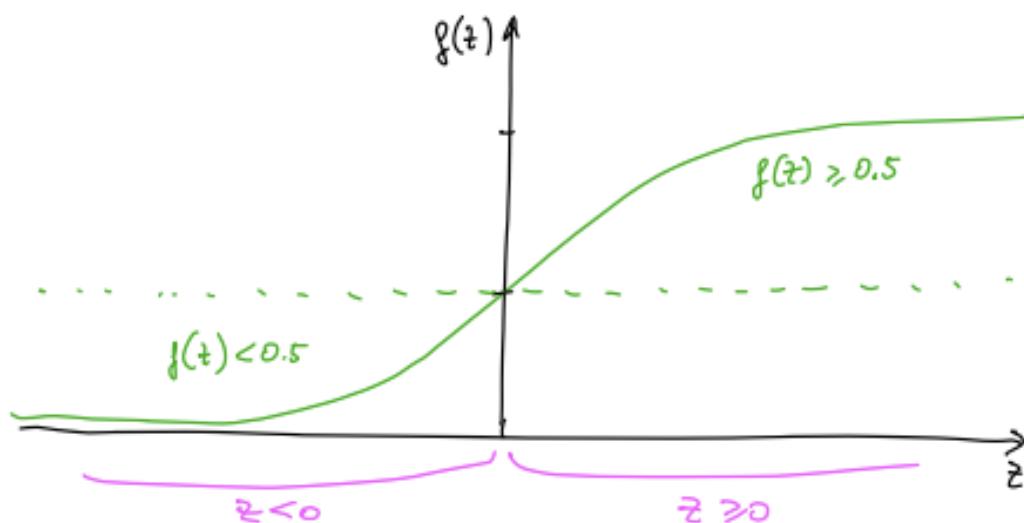
$$\Rightarrow P(y=0 | x; \theta) = 1 - P(y=1 | x; \theta)$$

Let's digest this..

$$h_{\theta}(x) = g(\theta^T x)$$

where $g(t) = \frac{1}{1+e^{-t}}$

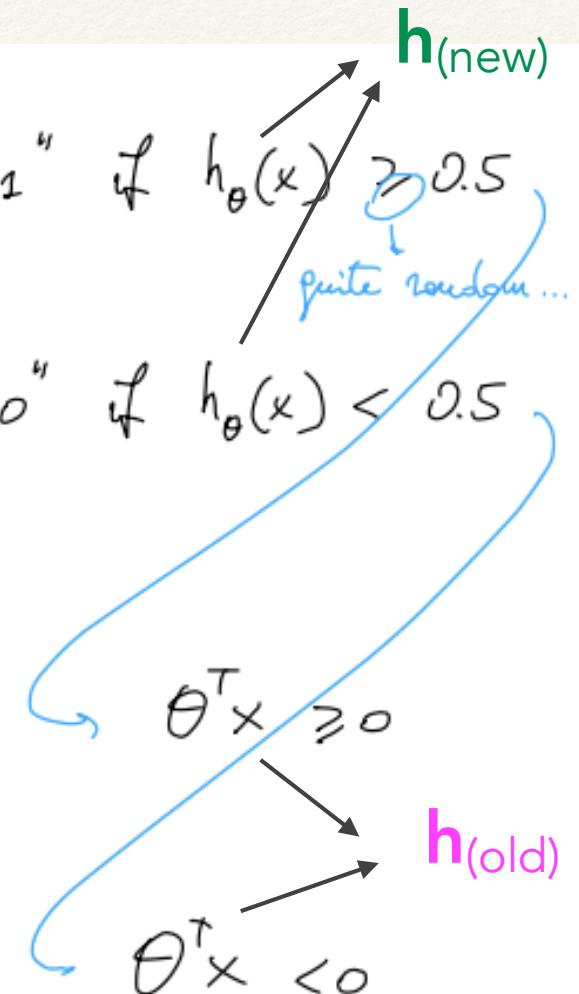
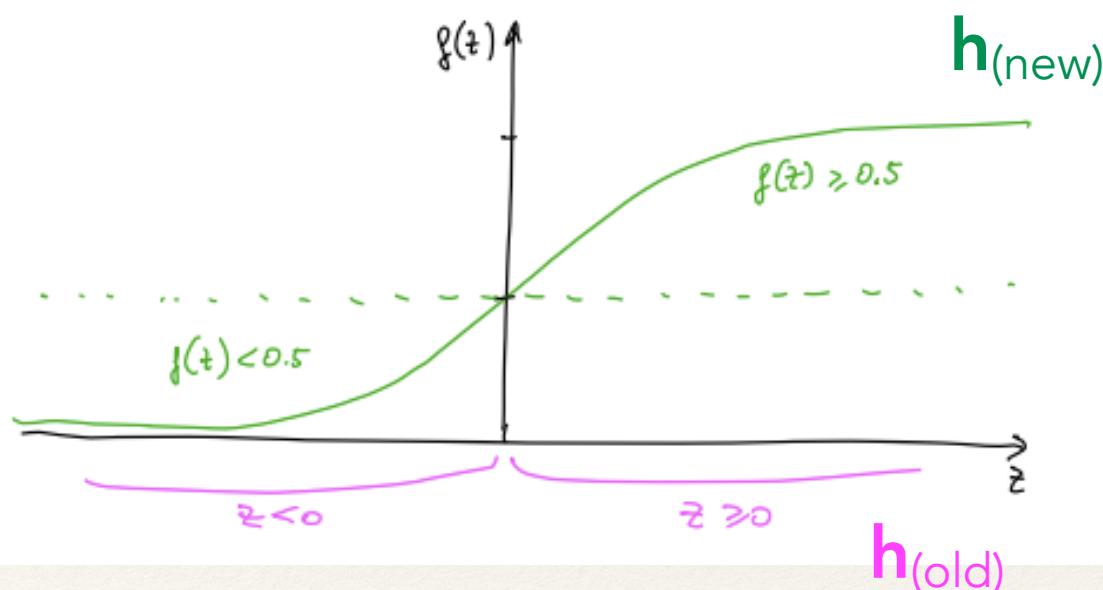
- predict " $y=1$ " if $h_{\theta}(x) \geq 0.5$
quite random...
- predict " $y=0$ " if $h_{\theta}(x) < 0.5$



From $\mathbf{h}_{\text{(old)}}$ to $\mathbf{h}_{\text{(new)}} = \mathbf{g}(\mathbf{h}_{\text{(old)}})$

$$\mathbf{h}_\theta(\mathbf{x}) = g(\boldsymbol{\theta}^\top \mathbf{x}) \quad \rightarrow \quad \text{where } g(z) = \frac{1}{1+e^{-z}}$$

- predict " $y=1$ " if $\mathbf{h}_\theta(\mathbf{x}) \geq 0.5$
quite random...
- predict " $y=0$ " if $\mathbf{h}_\theta(\mathbf{x}) < 0.5$



Summary

We changed the form for our hypotheses h in a way it is more solid against additional data points, and by the way it also resumes the fact it belongs to the $[0, 1]$ range

We did so by plugging the “old” h itself into the **logistic function** that becomes the “new” h (which we called $g(z)$).

The “new” h maps any real number into the $[0, 1]$ interval

- hence making it easy for transforming an arbitrary-valued function into a classification-friendly function

We should interpret h as the **probability** that our output is 1, with the prediction it is 0 being its complement.

1110001110100100110011010010100111010100100100101011
100101010010101010101101001010101011100011101001
00110011010010100111010010010101110010101001010
110110 Logistic Regression 010011010100101110101
00101010011010100110010010101110101010010101010
11010010101010111000111010010011001101001010011
0101101101011100101001010101010110100010101011
100011101001001100110100101001110101001001010110
0101010010011101011011011001010100101010101011
01000101110 Classification 10 and 10 representation 001
0111011100011101001001100110100101001110100100
10010010100 Decision boundary 00110010001000011111
10101101000101010111000111010010011001100010100
00100110011010010100011011110001001100110010100
111010101110010101010101110100110111010100111
010111011010101010101011001010001110000110001

Decision boundary

To get a better sense of what the logistic regressions hypothesis function is actually computing, we need to introduce the concept of **decision boundary**

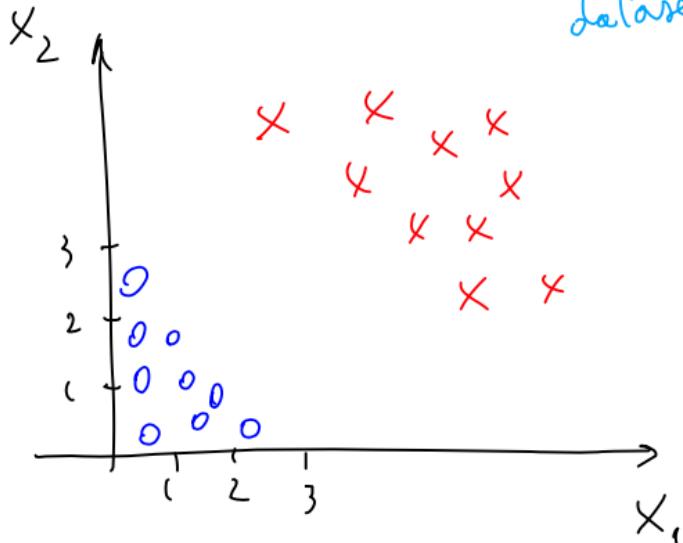
- we briefly introduce its concept, now we see it in some more details

We discuss:

- linear decision boundary
- non-linear decision boundary

How logistic regression makes its predictions

Example:



Suppose this dataset:

Suppose this hypothesis:

$$h_{\theta}(x) = f(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

Suppose we already fit the values of θ s:

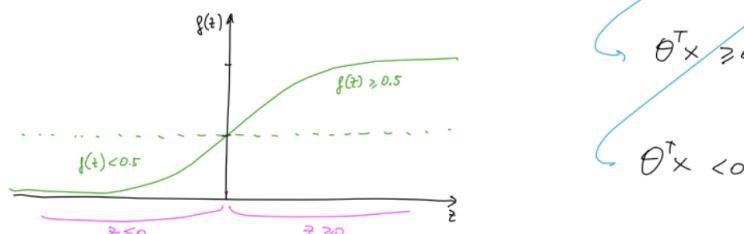
$$\begin{matrix} -3 \\ +1 \\ +1 \end{matrix} \Rightarrow \theta = \begin{pmatrix} -3 \\ +1 \\ +1 \end{pmatrix}$$

\times positive examples
 \circ negative examples

$h_{\theta}(x) = f(\theta^T x) \rightarrow$

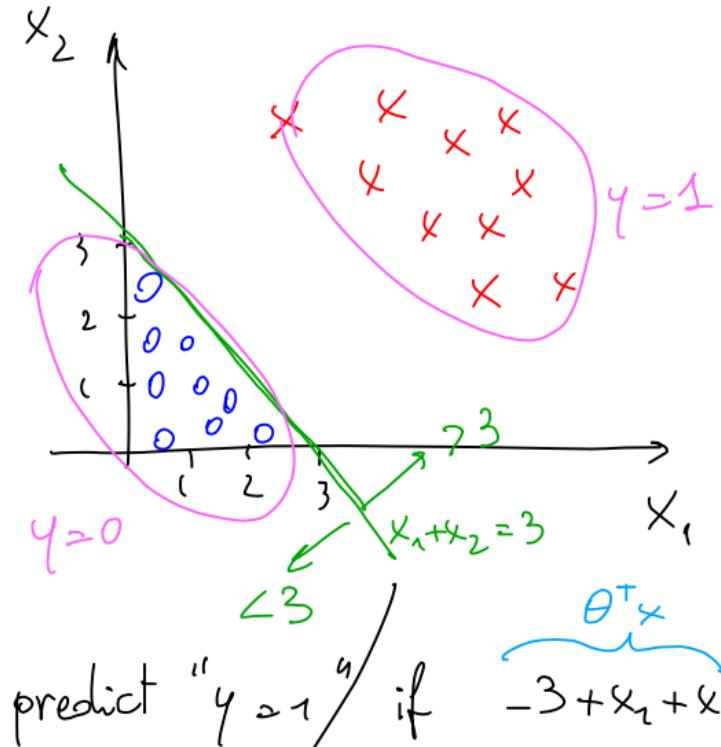
where $f(z) = \frac{1}{1+e^{-z}}$

- predict " $y=1$ " if $h_{\theta}(x) \geq 0.5$
quite random...
- predict " $y=0$ " if $h_{\theta}(x) < 0.5$



How logistic regression makes its predictions

Example :



Called DECISION BOUNDARY :

line that separates
region that predicts $y=0$ from
 $\nwarrow \uparrow \nwarrow \uparrow \nwarrow \uparrow$ $y=1$

Suppose this hypothesis :

$$h_{\theta}(x) = f(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

Suppose we already fit the values of θ s :

$$\begin{matrix} -3 \\ +1 \\ +1 \end{matrix} \Rightarrow \theta = \begin{pmatrix} -3 \\ +1 \\ +1 \end{pmatrix}$$

Decision boundary: a property of the hypothesis

So, the decision boundary is, in the feature space, the “border” between the regions where we predict $y = 1$ vs the regions where we predict $y = 0$

Note: **the decision boundary is a property of the hypothesis, including the parameters θ . It is not a property of the training set.**

- i.e. it is still of the same type (e.g. a line) even if e.g. we would change all data points in the training set
 - ❖ Later on we will talk about how to fit the θ s and there we'll end up using the training set, i.e. using our data: the data drives the training, not the type of decision boundary.

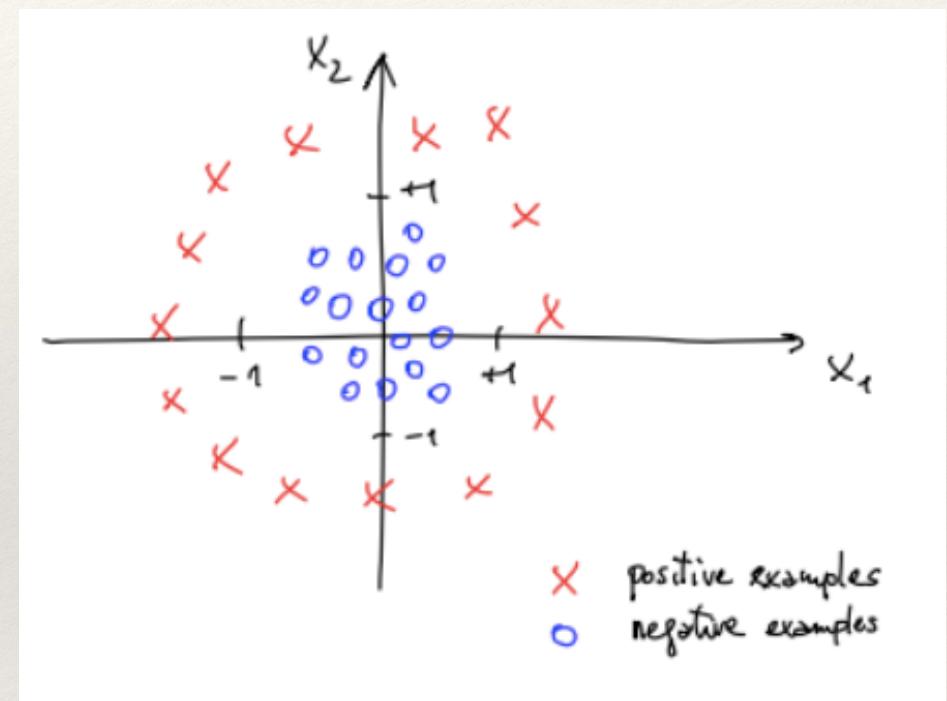
Recap:

- **the training set drives the training and it is used to fit the θ parameters...**
- ...but once you have the θ s for your hypothesis, it is **the hypothesis** that **defines the type of decision boundary**
 - ❖ in other words, you don't actually need to plot any training set in order to figure out the decision boundary - just study your hypothesis! In this way, a decision boundary linked to an hypothesis applies to other datasets too

Non-linear decision boundary

Let's now look at a more complex example.

Given a training set like the one on the right, how can I get logistic regression to fit this data?

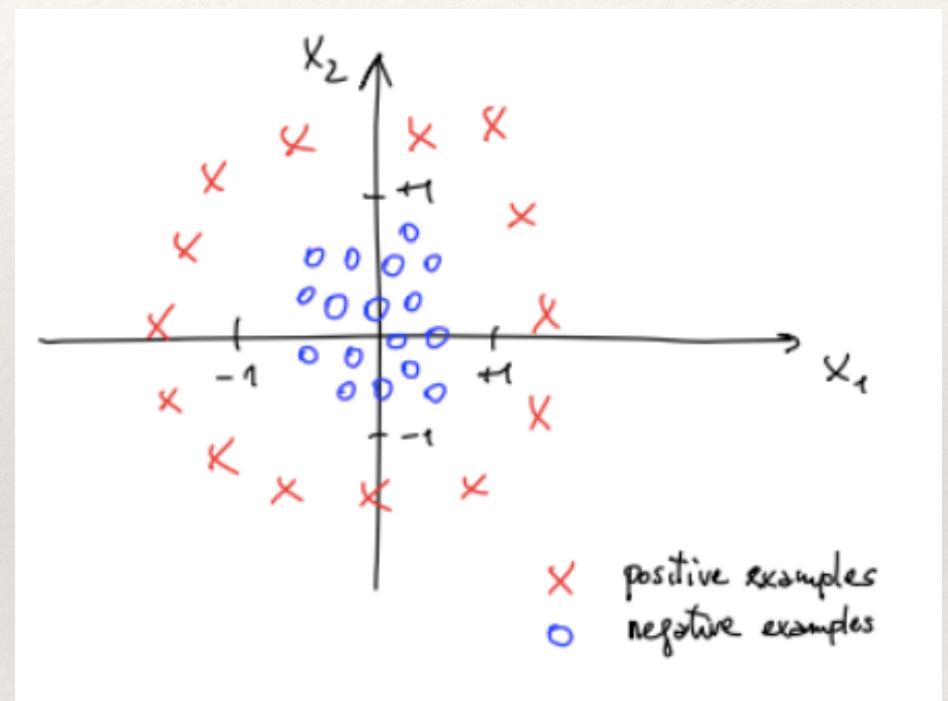


Non-linear decision boundary

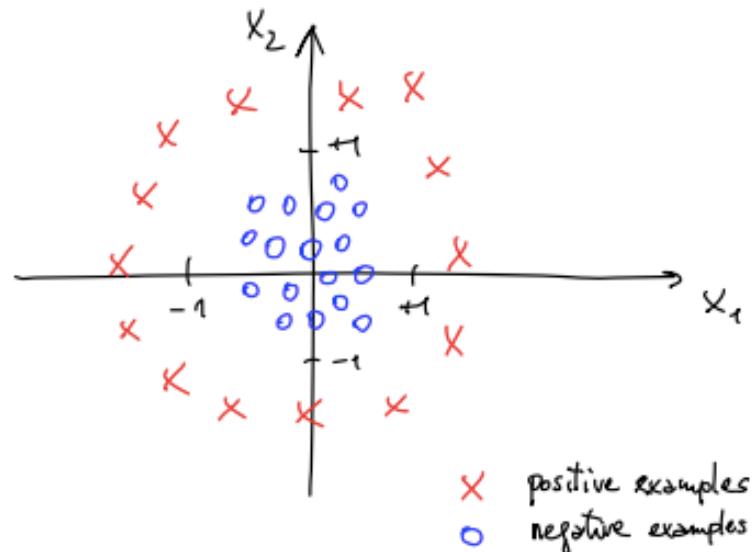
Let's now look at a more complex example.

Given a training set like the one on the right, how can I get logistic regression to fit this data?

Idea: in *regression*, we discussed about polynomial regression, i.e. adding extra higher order polynomial terms to the features vector. Why not doing the same for *logistic regression*?



Non-linear decision boundary



Suppose this hypothesis :

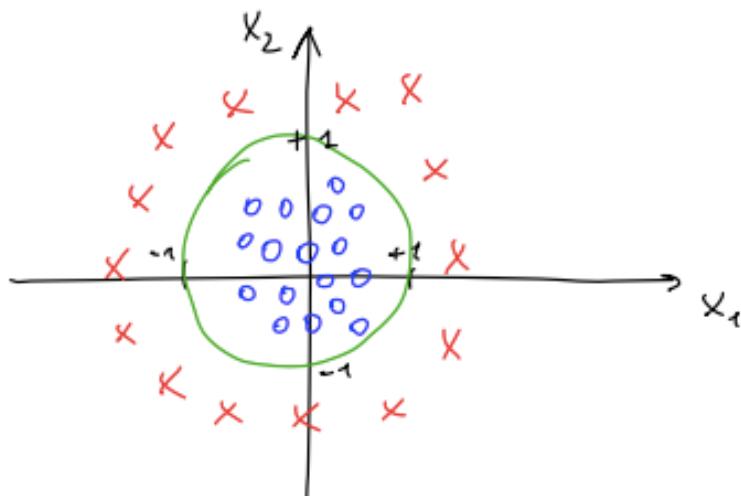
$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4)$$

Suppose we already fit the values of θ s :

$$\begin{matrix} -1 & 0 & 0 & 1 & 1 \\ \theta_0 & \theta_1 & \theta_2 & x_1^2 & x_2^2 \end{matrix}$$

$$\theta = \begin{pmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{pmatrix}$$

Non-linear decision boundary



Suppose this hypothesis :

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4)$$

θ_0 $\theta_1 x_1^2$ $\theta_2 x_2^2$ $\theta_3 x_3$ $\theta_4 x_4$

Suppose we already fit the values of θ s :

-1 0 0 1 1

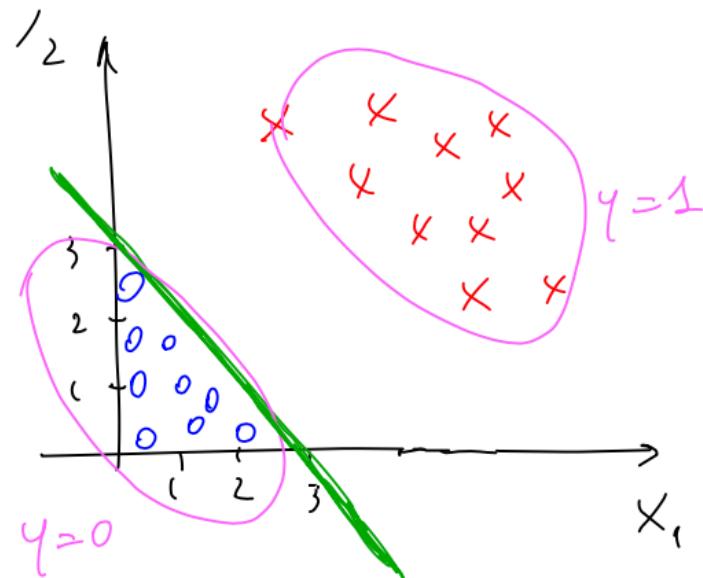
predict $y=1$ if $-1 + x_1^2 + x_2^2 \geq 0$
 0 < 0



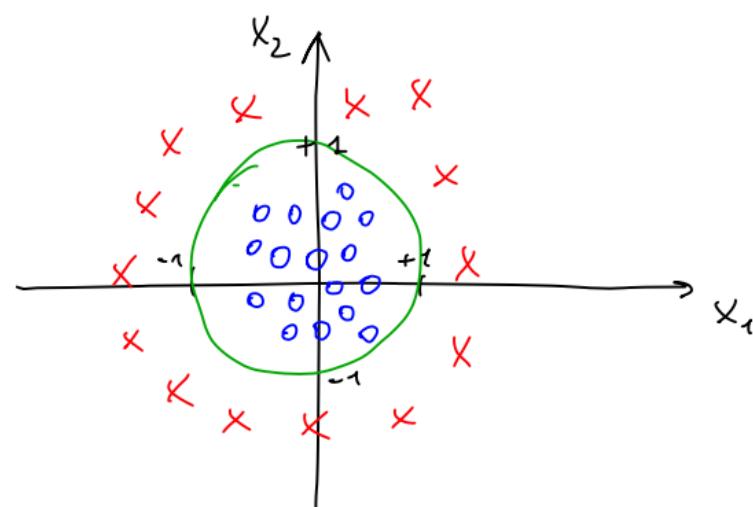
$$x_1^2 + x_2^2 = 1 \quad \text{is the decision boundary}$$

$$\theta = \begin{pmatrix} -1 \\ 0 \\ 0 \\ 1 \\ 1 \end{pmatrix}$$

Decision boundaries: property of.. ?



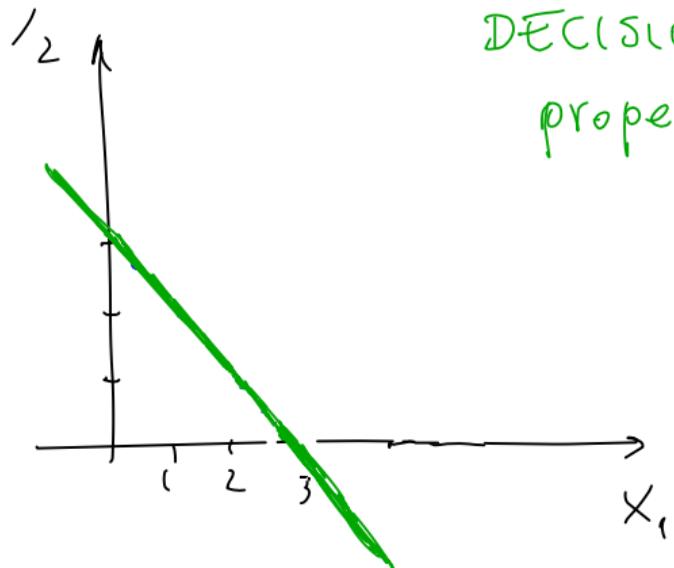
$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4)$$

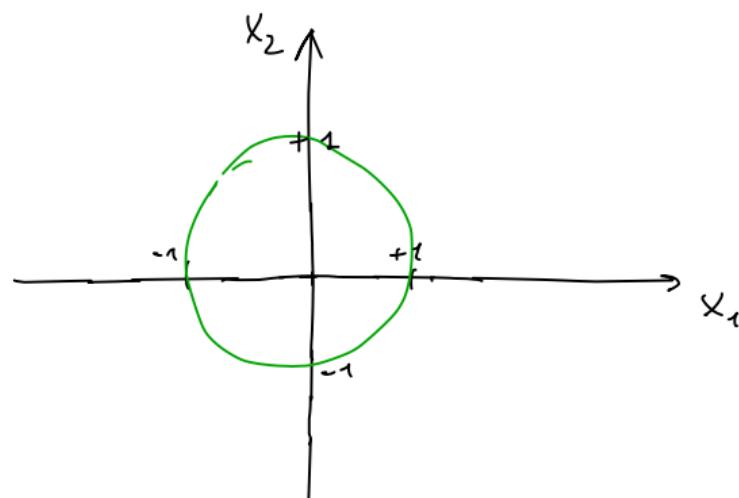
$x_1^L \quad x_2^L$

Decision boundary: a property of the h



DECISION BOUNDARIES as
properties of the hypothesis!

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$



$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \theta_4 x_4)$$

x_1^L x_2^L

Summary

In order to get our discrete 0 or 1 classification, we translated the output of the hypothesis **h** function into a logistic function **g**.

- When its input is greater than or equal to zero, its output is greater than or equal to 0.5.

This turns out to be as follows:

$$\theta^T x \geq 0 \Rightarrow y=1$$

$$\theta^T x < 0 \Rightarrow y=0$$

The **decision boundary** is the line that separates the area where $y = 0$ and where $y = 1$. **It is created by our hypothesis function.**

The input to the sigmoid function $g(z)$ (e.g. $\theta^T x$) doesn't need to be linear

- it could be a function that describes e.g. a circle (e.g. $z=\theta_0+\theta_1x_1^2+\theta_2x_2^2$) or any shape to fit your data

1110001110100100110011010010100111010100100100101011
100101010010101010101101001010101011100011101001
00110011010010100111010010010101110010101001010
111010 Logistic00 Regression01001101010010111010101
00101010011010100110010010101110101010010101010
11010010101010111000111010010011001101001010011
01011011010111001010010101010101101000101010111
1000111010010011001101001010011101010010010101110
0101010010011101011011011001010100101010101011
010001010110 Logistic10 regression10 model0100100100
1011101111000111010010011001101001010011101010010
010010010100 Cost10 function00110010001000011111101
011010001010101110001110100100110011000101000001
001100110100101000110111100010011001100100100111
01010111001010101010101110011001101110101010011101
1110110101010101010101100101010001110000110001101

Cost function for logistic regression

Let's discuss how to fit the parameters θ for logistic regression.

In particular, let's define the optimisation objective, i.e. the **cost function** that we will use to fit the parameters.

Fitting a logistic regression model

training set

$$\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$$

m examples

m examples

vectors

$$x \in \mathbb{R}^{n+1} \begin{pmatrix} x_0 = 1 \\ x_1 \\ \vdots \\ x_n \end{pmatrix}$$

$$y \in \{0, 1\}$$

classification

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

hypothesis parameter

How to choose parameters θ ?

Building a cost function for logistic regression

Let's rewrite the cost function for LINEAR REGRESSION :

$$\begin{aligned} J(\theta) &= \frac{1}{2m} \sum_{i=1}^m \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2 = \\ &= \frac{1}{m} \sum_{i=1}^m \underbrace{\left[\frac{1}{2} \left(h_\theta(x^{(i)}) - y^{(i)} \right)^2 \right]}_{\text{"cost } (h_\theta(x^{(i)}), y^{(i)})"} = \\ &= \frac{1}{m} \sum_{i=1}^m \text{cost } (h_\theta(x^{(i)}), y^{(i)}) \end{aligned}$$

sum over my
training set

For linear regression this works just fine. But...

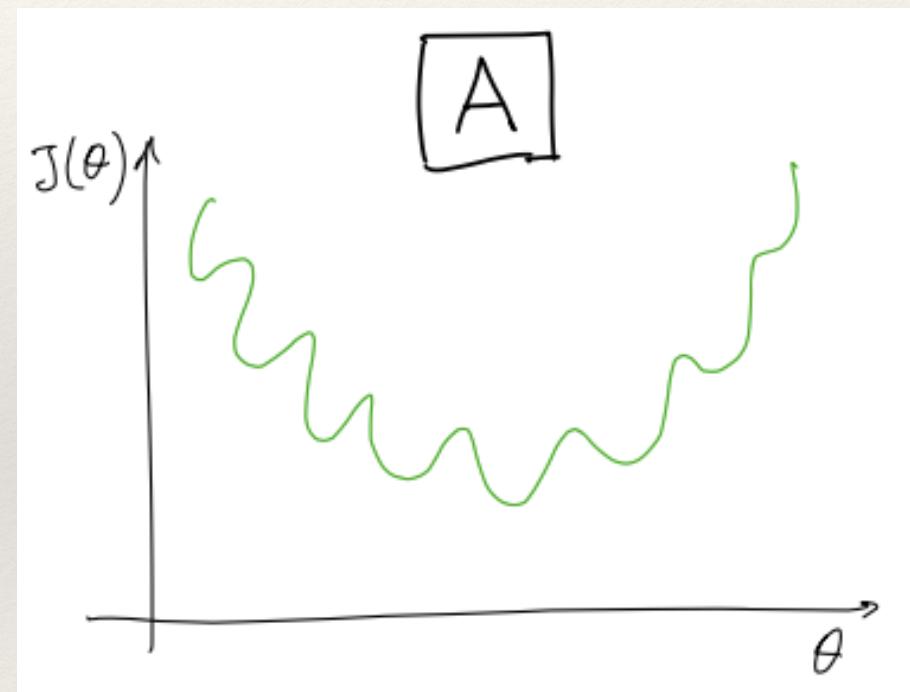
Building a cost function for logistic regression

..what if we try exactly this for logistic regression?

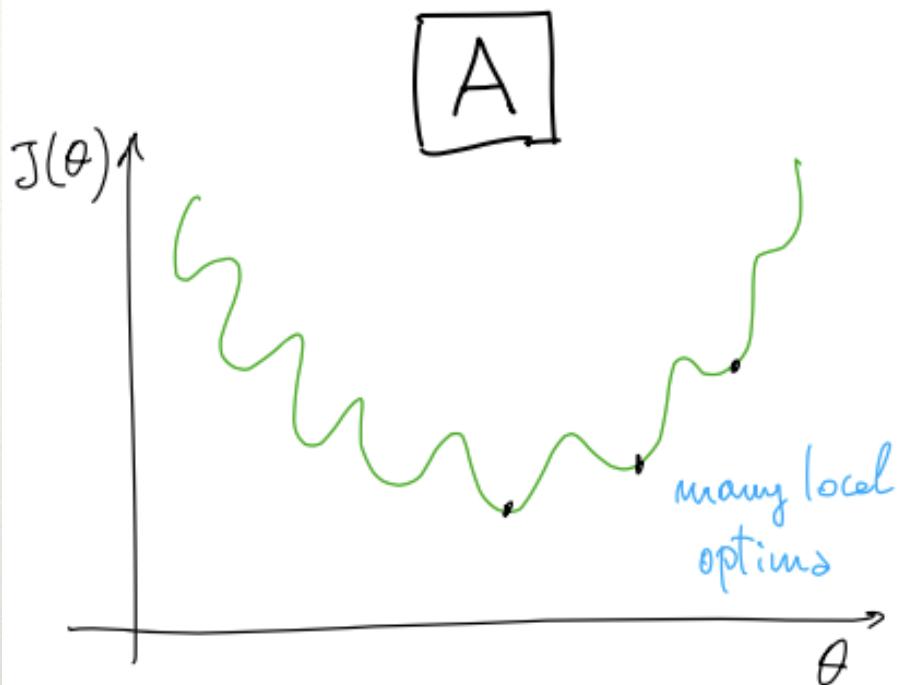
- in principle, if we could minimise this “cost()” part that is plugged into J , that would still work OK also for logistic regression.

But it turns out that - for logistic regression - we can't.

- If we use this particular cost function, this would be a non-convex function of the parameters.
 - ❖ this, as a candidate cost function $J(\theta)$ for logistic regression too, is built on a hypothesis h that has a strong nonlinearity (it is the sigmoid function!). So if you take that function, plug it in the “cost()” part and plug that into the $J(\theta)$, and then plot $J(\theta)$, it may look like plot A on the right
- if you were to run GD on this sort of function it is not guaranteed *at all* to converge to the global minimum!



We are stuck with this..



$h_\theta(x)$ non linear

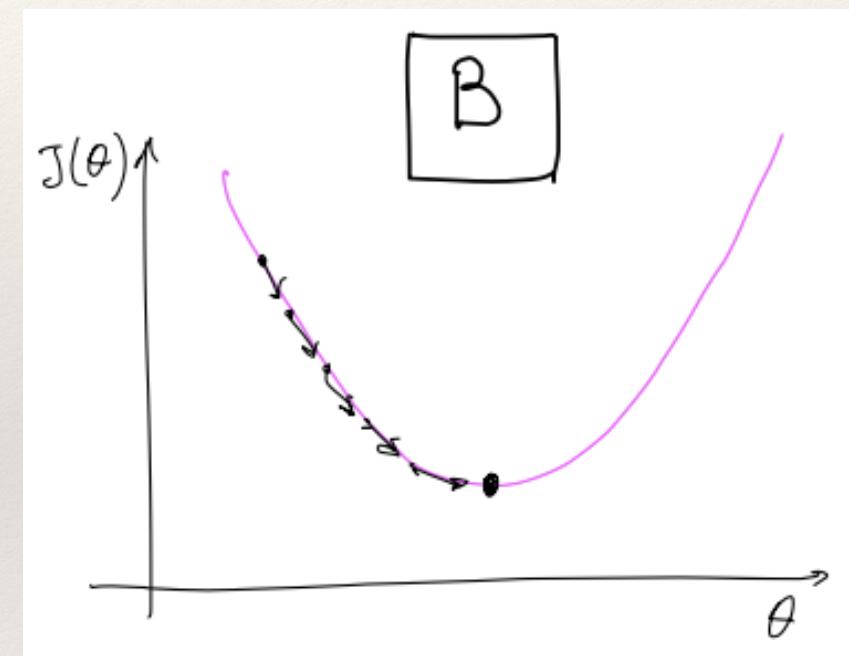


$J(\theta)$ non convex

Building a cost function for logistic regression

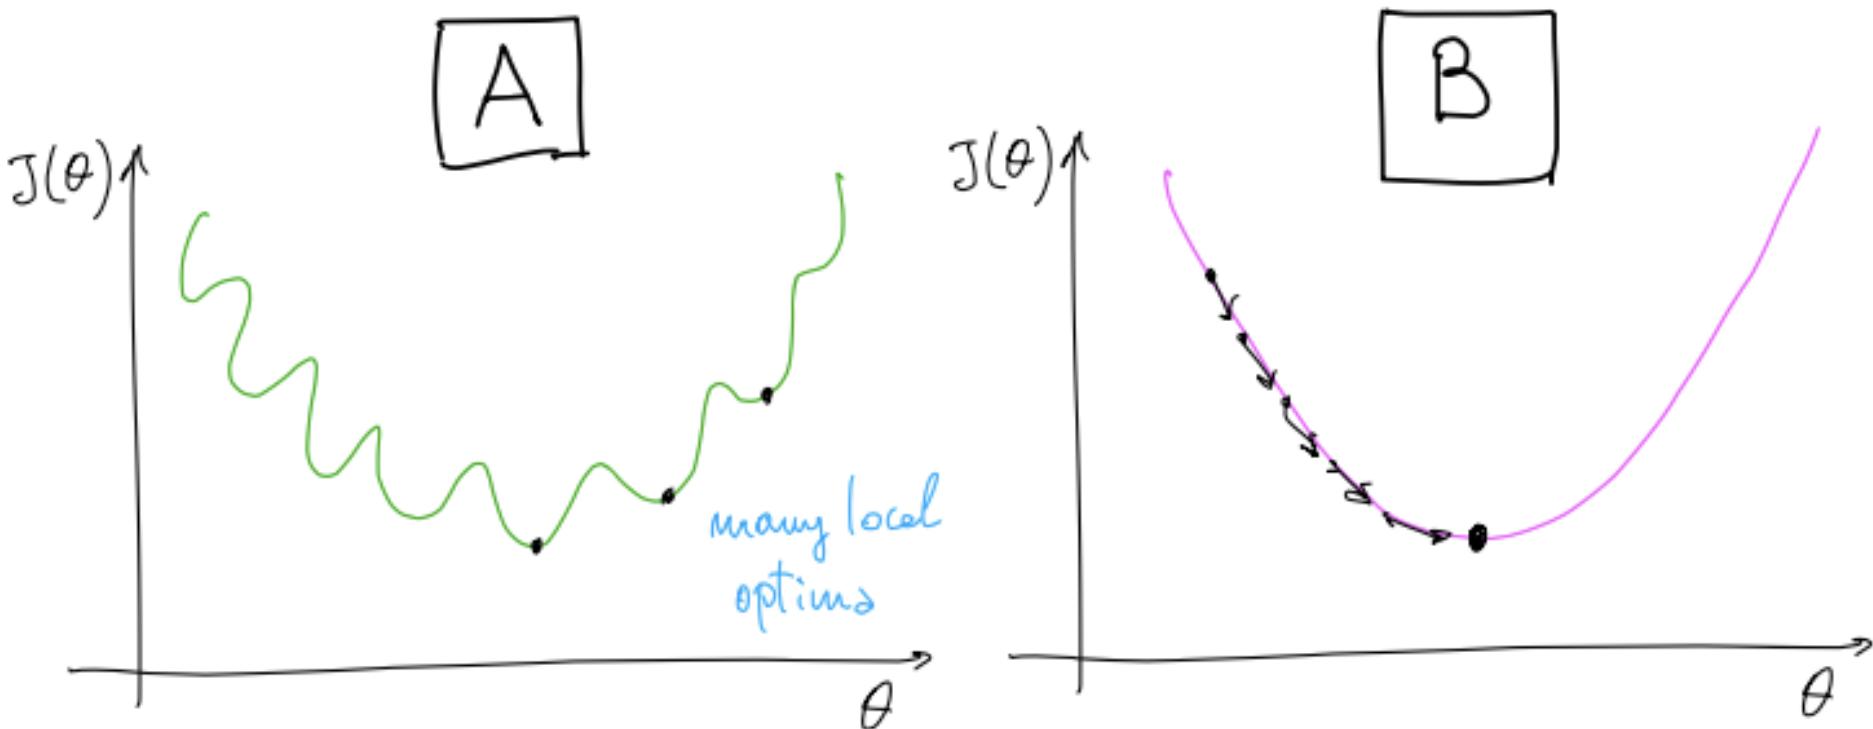
Instead, what we would like is to have a cost function $J(\theta)$ that is convex ("bowl-shaped") - see B - so that we would be guaranteed that GD would converge to a (possibly global) minimum.

Well, nobody said - given the un practicality of having a very nonlinear sigmoid function that forces $J(\theta)$ to be non-convex - that we need to define a cost function indeed as a squared cost function!



So: let's come up with something that allows $J(\theta)$ to be convex, so we can apply GD and we are guaranteed to find a good (global) minimum..

Pictorial recap on building a cost function for logistic regression



$h_\theta(x)$ non linear



$J(\theta)$ non convex

which $h_\theta(x)$???

~~which $h_\theta(x)$???~~

$J(\theta)$ CONVEX

Logistic regression cost function

[Note: we drop " (i) "]

$$\text{cost} (h_{\theta}(x), y) = \begin{cases} -\log (h_{\theta}(x)) & \text{if } y=1 \\ -\log (1-h_{\theta}(x)) & \text{if } y=0 \end{cases}$$

NOTE: Intuition goes first. The topic of convexity analysis is beyond the scope of this course, but it is possible to show that with this particular choice of cost function, this will give a convex optimisation problem, i.e. $J(\theta)$ will be convex and local optima free.

Remember:

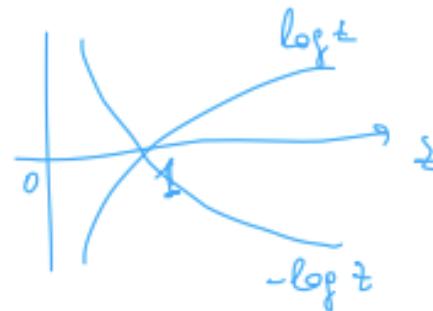
$$\begin{aligned} J(\theta) &= \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \\ &= \frac{1}{m} \sum_{i=1}^m \underbrace{\frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2}_{\text{"cost } (h_{\theta}(x^{(i)}), y^{(i)})"} = \\ &= \frac{1}{m} \sum_{i=1}^m \text{cost } (h_{\theta}(x^{(i)}), y^{(i)}) \end{aligned}$$

sum over my
training set

Logistic regression cost function

[Note: we drop " (i) "]

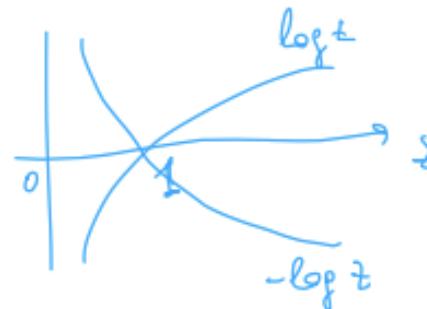
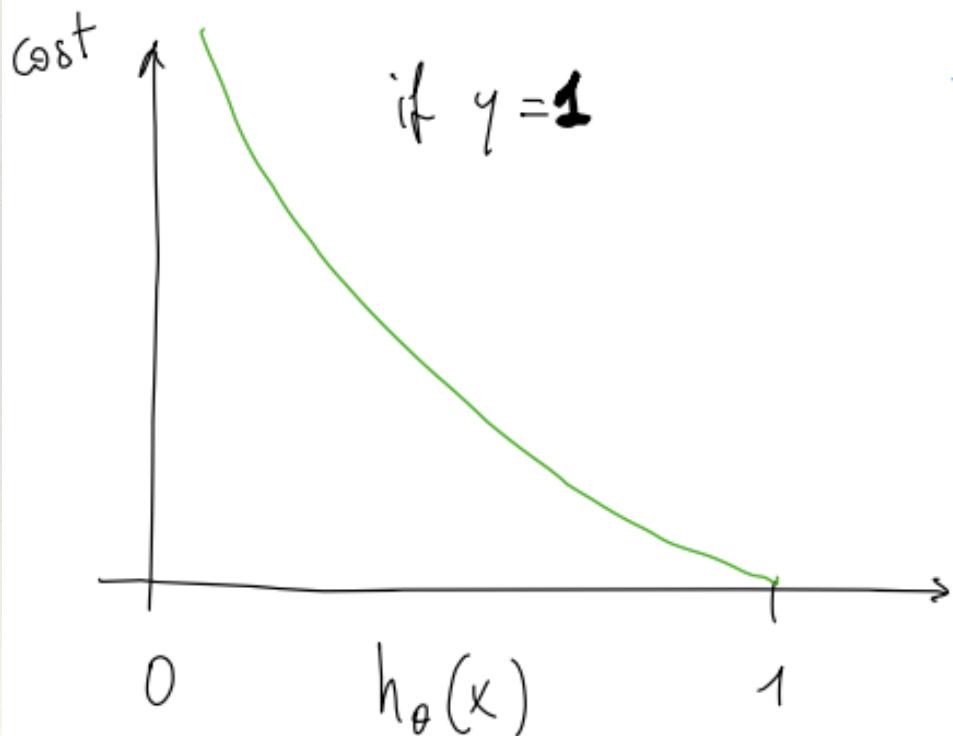
$$\text{cost } (h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y=1 \\ -\log(1-h_\theta(x)) & \text{if } y=0 \end{cases}$$



Logistic regression cost function ($y=1$)

[Note: we drop " (i) "]

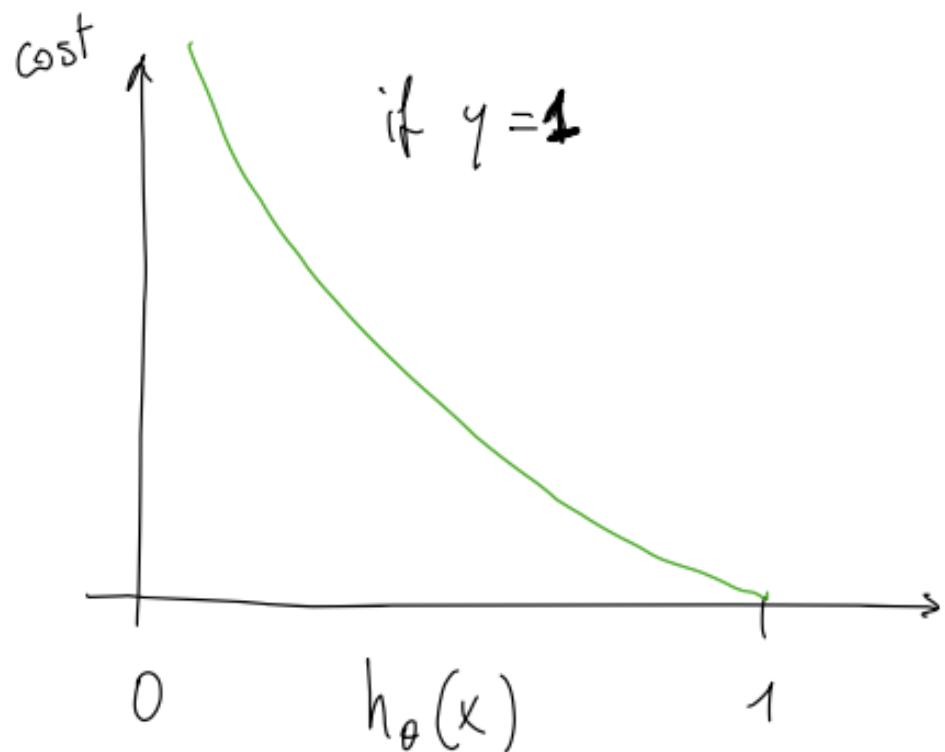
$$\text{cost} (h_\theta(x), y) = \begin{cases} -\log(h_\theta(x)) & \text{if } y=1 \\ -\log(1-h_\theta(x)) & \text{if } y=0 \end{cases}$$



Logistic regression cost function ($y=1$)

[Note: we drop " (i) "]

$$\text{cost} (h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1-h_{\theta}(x)) & \text{if } y=0 \end{cases}$$



Nice! Because:

$$h_{\theta}(x)=1, y=1 \Rightarrow \text{cost} = 0$$

$$\text{as } h_{\theta}(x) \rightarrow 0, y=1 \Rightarrow \text{cost} \rightarrow \infty$$

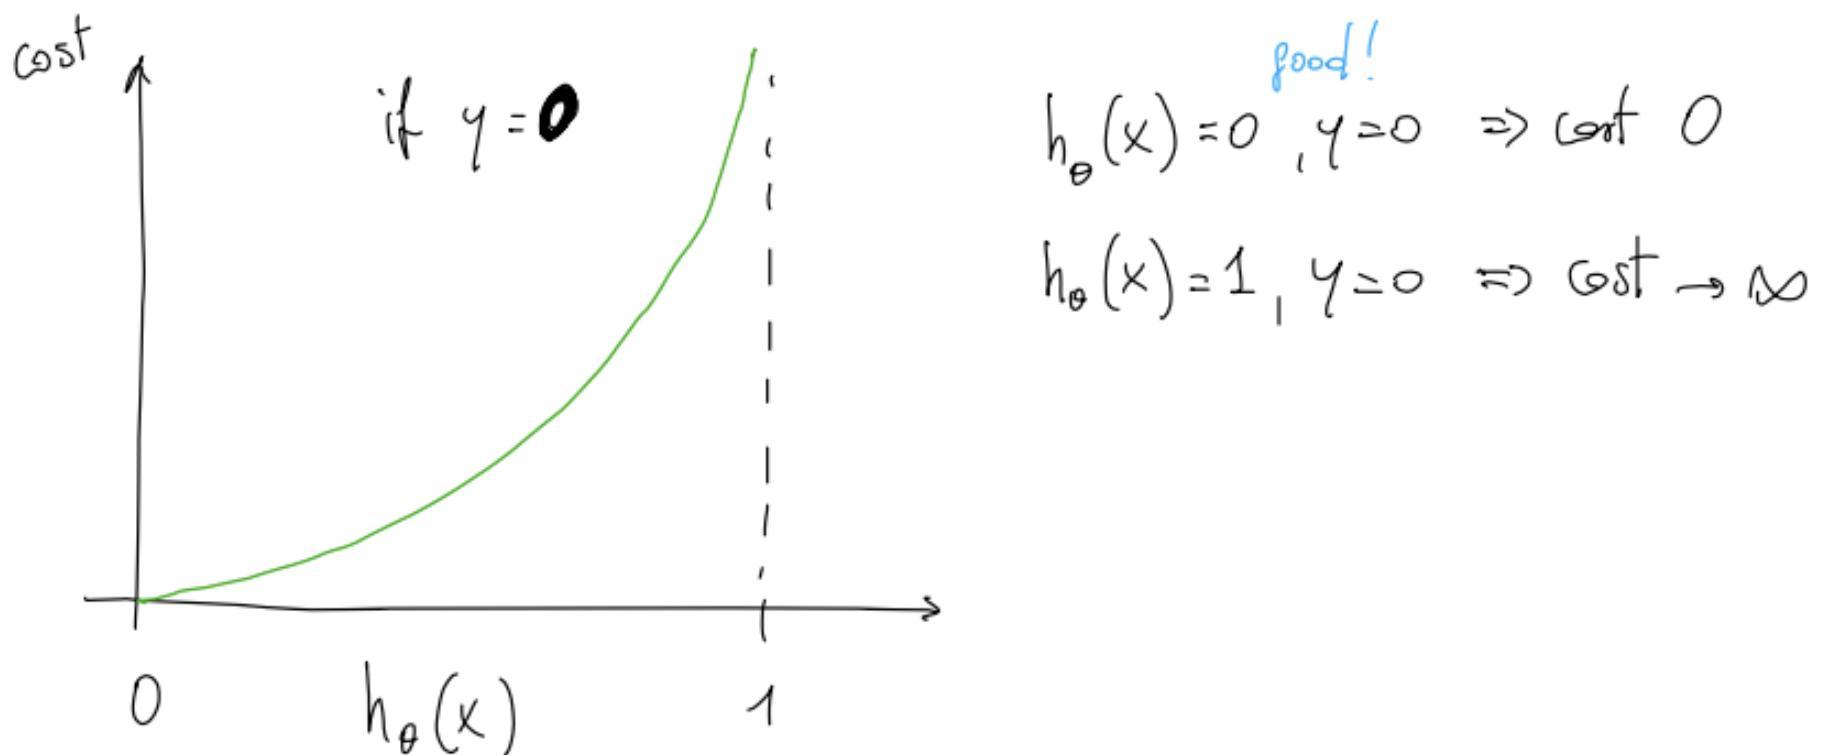
Able to capture intuition that if $h_{\theta}(x)=0$ (predict $P(y=1|x;\theta)=0$), but $y=1$ \Rightarrow it penalizes learning algo by a very large cost!

think of tumors!

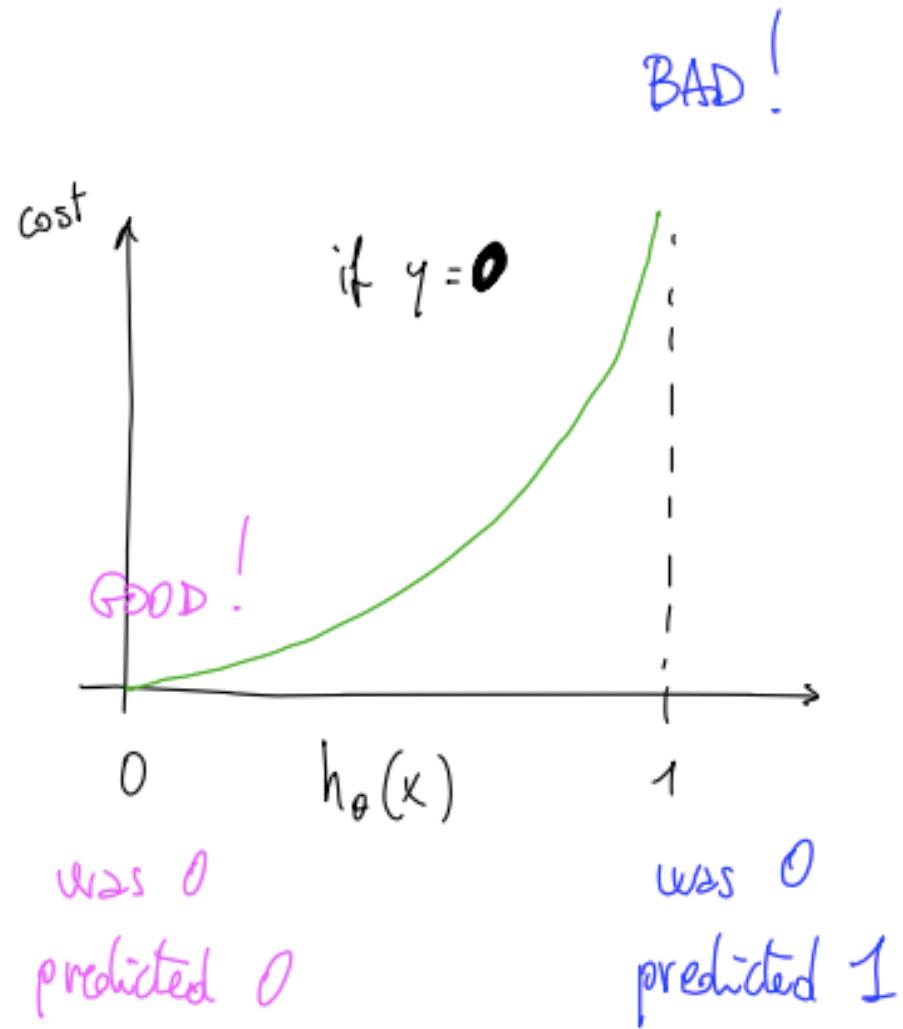
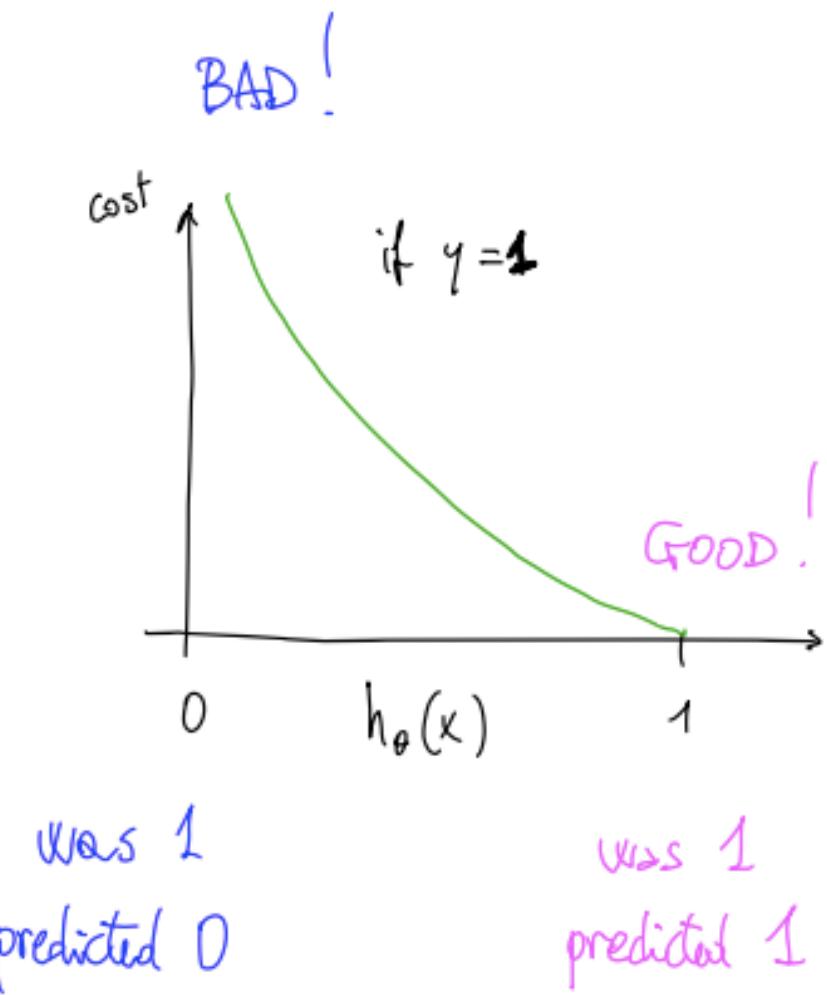
Logistic regression cost function ($y=0$)

[NOTE: we drop " (i) "]

$$\text{cost} (h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1-h_{\theta}(x)) & \text{if } y=0 \end{cases}$$



Recap on logistic regression cost function



Summary

We cannot use the same cost function used for linear regression because the logistic function will cause the output to be wavy, causing many local optima.

We then managed for logistic regression to define a cost function that is convex.

[Note : we drop " (i) "]

$$\text{cost} \left(h_{\theta}(x), y \right) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y=1 \\ -\log(1-h_{\theta}(x)) & \text{if } y=0 \end{cases}$$

But, we are not yet there..

So far, we have defined the cost function for a single training example. Now we are going to develop this further, and define the cost function for the entire training set. This will bring us to a simpler way to write it, and then we will work out GD for logistic regression.